Differing Priors and Information Acquisition When Others Make the Decision*

Dario Trujano-Ochoa†

October, 2026

Link to the Latest Version

Abstract

Should you pay more for information when someone else will act on it instead of you? Standard Bayesian theory predicts it should not matter, and this paper tests that prediction experimentally.

Pairs face the Bayesian decision problem of purchasing a signal about a binary state. I elicit first- and second-order posterior beliefs to evaluate the effects of different priors on belief updating. Then, I independently vary (i) whether participants share common or differing prior beliefs when eliciting their posteriors, and (ii) whether the information purchaser is also the final decision-maker.

Three main results emerge. First, participants expect others to update beliefs similarly to themselves. Second, sensitivity to information predicts willingness to pay (WTP): participants who put more weight on the signal bid more to implement their own strategy. Third, participants exposed to differing priors bid more to implement their own strategy than the other participant's, a premium absent with common priors.

This premium is not explained by the instrumental value of the strategy, which is consistent with a non-instrumental value of decision rights that becomes salient when beliefs may differ. This has implications for organizational information acquisition and delegation design.

JEL Classification: D83, D91, C91

Keywords: Information Acquisition, Delegation, Higher-Order Beliefs, Differing Priors

*I want to thank Gary Charness, Erik Eyster, Daniel Martin, Ryan Oprea, Sevgi Yuksel, and all the people who discussed with me the early ideas and versions of this project at UCSB. The paper was improved at Texas A&M, where I received feedback and comments from Marco Palma, Paul Feldman, Marissa Lepper, Catherine Eckel, and the collaborators at the Human Behavior Laboratory (HBL). This experiment was funded by the Department of Economics at UCSB and completed while the author was a Postdoctoral Researcher at the HBL at Texas A&M University. This material is based upon work supported by the Texas A&M AgriLife Institute for Advancing Health Through Agriculture, IHA, and the U.S. Department of Agriculture, Agricultural Research Service, under Agreement No. 59-3093-4-00. Any opinions, findings, conclusions, or recommendations expressed in this publication are those of the author(s) and do not necessarily reflect the view of the U.S. Department of Agriculture.

†Texas A&M University. dario.trujanoochoa@tamu.edu

1 Introduction

In many real-world settings, the individual who decides whether to acquire information is not the same as the one who ultimately takes the action based on that information. This separation between information acquisition and decision-making is common in organizations and policy environments. In such settings, the value of information depends not only on its statistical properties but also on how it is expected to be used by others.

Additionally, information acquisition and final decisions can be made while in a context of interacting with others who hold the same or different beliefs. Such differences in prior beliefs can be found in political opinions, investment decisions, and organizational choices. These differences could shape how individuals interpret new information and how much they value acquiring additional information.

Motivating example. Organizations frequently face decisions under uncertainty and may reduce risk by collecting additional information, such as through sampling or expert advice. However, acquiring information is costly, and it is not always optimal to do so. Moreover, the decision to acquire information and the final decision based on that information are not necessarily made by the same individual. To illustrate this, consider the following example. Anne and Ben must plan inventory for multiple products, where demand for each product is either high or low, and both states are initially considered equally likely. Before placing orders, they can hire a consultant to provide a demand forecast for a fixed fee. In an 'integrated' decision structure, Anne and Ben are each responsible for different products and independently decide whether to purchase the consultant's forecast and how much inventory to order. In a 'split' decision structure, Anne determines the inventory levels for all products, while Ben decides whether to incur the cost of the consultant's forecast. This distinction captures a key organizational feature: the separation between information acquisition and action.

Standard Bayesian models with aligned incentives imply that the value of information is independent of who makes the final decision. If all agents update beliefs correctly and share the same objective, information should be equally valuable regardless of whether it is used by oneself or by another decision-maker. Then Ben should be willing to pay the same amount for the consultant under both the integrated and split decision structures. This prediction relies on the assumption that individuals expect others to process information in the same way they do.

This paper tests this implication in an experimental setting where participants may hold differing prior beliefs and form expectations about how others update their beliefs. I study two related questions. First, how do individuals update their own beliefs, and how do they expect others to update when priors differ? Second, how do these beliefs and expectations affect their willingness to pay (WTP) for information when decisions are implemented either by themselves or by another participant?

To address these questions, a laboratory experiment was designed in which participants observe signals about a binary state and report posterior beliefs for themselves and for another participant. The experiment manipulates whether participants share a common prior (Common Priors, CP, treatment) or report the posterior of someone whose prior differs from their own (Differing Priors, DP, treatment); treatments were assigned by session. In a subsequent stage, participants state their WTP for acquiring information under two decision structures: an integrated structure, where they implement their own strategy, and a split structure, where another participant makes the final decision.

The results show that, on average, participants expect others to update beliefs in a manner similar to their own. In the DP treatment, there was a larger variance of the difference between individuals' own updating strategies and their expectations of others, but the mean was not different from zero.

WTP for information was highest when participants in the DP treatment implemented their own strategy (\$0.82, compared with \$0.64 when they implemented the other participant's strategy; $p = 0.03$). In the CP treatment, the two decision structures produced similar bids (\$0.66 and \$0.62; $p = 0.65$). Bids for the other participant's strategy are similar in both treatments, so the difference comes from bids for one's own strategy. Because treatments were assigned by session, I also test the treatment differences by re-assigning treatments across sessions (randomization inference). With this test, the difference between the two gaps is not statistically significant ($p = 0.29$), and own-strategy bids are higher in the DP sessions with $p = 0.10$ ($p = 0.05$ when the comparison includes every session without differing priors). The pattern is therefore consistent with exposure to differing priors raising the value of implementing one's own strategy, but it does not establish it. Because the strategies are fixed in Section 1, this premium does not come from the opportunity to revise the decision.

To better understand these patterns, I measure individual deviations from Bayesian updating, focusing on conservatism and base-rate neglect. Sensitivity to information

predicts WTP: participants who put more weight on the signal bid more to implement their own strategy, and the value of the signal implied by their own updating predicts these bids. Moreover, individuals' expectations of others' updating are correlated with their own biases, suggesting that participants project their own updating behavior onto others.

Two sets of additional sessions support these conclusions. When each participant reported either their own posteriors or a prediction of someone else's, but not both, the predictions were still indistinguishable from own posteriors. When strategy tables that were not monotone in the signal were flagged, updating moved closer to Bayes' rule and participants copied their own posteriors less often when predicting others, but WTP did not change, so these sessions are pooled with the main sample in the WTP analysis. When WTP was elicited at priors of 20%, 50% and 90%, bids did not respond to the prior, even though the Bayesian value of the signal varies sharply with it.

These findings show that sensitivity to information predicts WTP, but also that two behavioral patterns should be considered. First, participants expect others to update like themselves, so implementing one's own strategy should add little instrumental value; yet with differing priors they bid more to implement it. This is consistent with an increase in the non-instrumental value of decision rights (Bartling et al., 2014; Owens et al., 2014) when participants are made aware that others may hold different beliefs. Second, WTP is not sensitive to the prior, as theory predicts. More broadly, the results suggest that the value of information in strategic settings depends not only on its statistical content but also on expectations about others' behavior.

This paper contributes to the literature on belief updating and information acquisition by showing that the valuation of information is sensitive to the interaction between differing priors and the allocation of decision rights. It also contributes to the growing literature on the behavioral foundations of information demand by documenting systematic deviations from the standard Bayesian benchmark in environments with aligned incentives but decentralized decision-making.

2 Related Literature

The literature on (non-) Bayesian updating is considerable, and the general results are that people display base-rate neglect and conservatism (Benjamin, 2019).[1] I know of no paper that explores whether people expect others to display the same biases as themselves when the priors differ or are the same.

This paper innovates in the elicitation of the willingness to pay (through the Becker–DeGroot–Marschak (BDM) mechanism) in a strategic setting where the level of bias expected from another person can be used to predict the value of providing information to this person. I also analyze whether differences in prior beliefs affect the willingness to pay to receive or send an informative signal.

2.1 Non-Bayesian Belief Updating

It has been consistently found that participants deviate from Bayes' rule; conservatism and base-rate neglect (BRN) are general features of belief updating. Benjamin (2019) presents a comprehensive analysis of the results in the laboratory. He used the model from Grether (1980, 1992) to estimate these biases. To predict the elicited posteriors in this paper, I denote with β^p and β^l the relative importance of the priors and the likelihood of the signal, respectively.[2]

2.2 Higher-Order Belief Updating

Whether people expect others to display the same belief-updating biases as themselves is understudied. The paper "Higher-Order Learning" by Evdokimov and Garfagnini (2022) was the first experimental effort to study how higher-order beliefs are updated. They elicited the higher-order beliefs of others' posterior beliefs in a sequential Bayesian task. These authors estimated the biases using the econometric approach of Grether (1980, 1992) (as in Benjamin 2019). They found considerable individual heterogeneity. However, the first- and higher-order beliefs parameter distributions were significantly different only at 10%,[3] and non-statistically different when the beliefs were elicited simultaneously on MTurk. Finally, they found a lack of conver-

[1] However, these biases are identifiable under different prior-likelihood combinations. For example, with a uniform (50–50) prior, only conservatism can be measured.

[2] Benjamin estimated $\beta_i^l = 0.383$ and $\beta_i^p = 0.434$ from a meta-analysis of independent experiments in settings similar to the design implemented in this study.

[3] The distribution of β^p was different between the first-order beliefs and second- and third-order beliefs, and β^l was different between the first- and second-order beliefs.

gence simulating data with the parameters estimated in the experiment. This was a property of Grether's quasi-Bayesian model with BRN also explored theoretically by Benjamin et al. (2019).[4] It was also found that participants with a higher cognitive ability expect others to end up agreeing more with them (Evdokimov and Garfagnini, 2023). Other studies have focused on higher-order beliefs on actions, risk preferences, and time discounting (Fedyk, 2021). Another important related work is from Agranov and Detkova (2024). They also studied expected belief updating from others and analyzed the effects of different priors, but they used statements. In that paper, there is an asymmetry in the expected belief of others after the signal, depending on participants having different priors. In this paper, I also analyze the willingness to pay to implement the strategy of another person.

It has been reported that people are blind to their own biases and frequently consider that others have more significant biases than themselves (Pronin et al., 2002; Pronin, 2008). However, the results from Evdokimov and Garfagnini (2022) show that participants in belief-updating tasks could be more prone to false consensus biases (Engelmann and Strobel, 2000). In the case of higher-order beliefs, people might report the same beliefs for others as for themselves, especially if the other is perceived as similar to you.

2.3 Strategic Settings

Other studies have considered higher-order beliefs in strategic settings. Szkup and Trevino (2020, 2026) applied a game where it is payoff-relevant to estimate the actual state of the world but also to match the guess from others. They elicited participants' beliefs about their opponents' decisions after an informative signal about the game's fundamentals. This task requires calculating the expected signals observed by the opponent and knowing how those signals were updated. They found that participants are accurate in their first-order beliefs, but they deviate from the theoretical predictions in their stated beliefs about the actions of others. Another closely related experiment is the one designed by Campos-Mercade and Mengel (2024). They considered studying the effects of higher-order beliefs on the conservatism of employers on employees' decisions. Participants in the role of employees were sensitive to heterogeneous biases, but this came exogenously as the experimenters implemented it, and they received feedback on the decisions made by the participants in the role of em-

[4]Intuitively, when agents base-rate-neglect, new signals move the posterior more than Bayes' rule implies, and the beliefs don't settle.

ployers. Because employers were classified according to their degree of conservatism in belief updating and then paired with employees, the latter had to learn the biases of employers. This setting is different from the present study since the question here is about the biases that participants have about the biases of others without getting any feedback.

2.4 Control Premium

This paper contributes to the literature on decision rights, delegation, and belief formation by examining how the value of information depends on who ultimately makes the decision. In particular, it relates to two strands of research: (i) the intrinsic value of decision rights, and (ii) delegation decisions shaped by beliefs about others.

2.4.1 Intrinsic Value of Decision Rights

A large literature documents that individuals value decision rights beyond their instrumental impact on outcomes. Owens, Grossman, and Fackler (2014) show that subjects are willing to pay to retain control even when delegating decisions would increase expected monetary payoffs. Similarly, Bartling, Fehr, and Herz (2014) provide evidence that individuals derive utility from having decision authority itself, linking control to autonomy and psychological ownership.

These findings suggest that preferences for control can generate a wedge between optimal and observed behavior. In standard models, the value of information depends solely on its effect on expected outcomes. However, if individuals intrinsically value making their own decisions, they may be willing to pay more for information when they retain control.

The present paper complements this literature by isolating a distinct mechanism. Rather than focusing on an intrinsic utility from control, I consider environments with aligned monetary incentives where control should not matter under Bayesian rationality. I show that differences in willingness to pay for information across decision structures can arise even in the absence of intrinsic preferences for control. Instead, these differences emerge from differing priors and expectations about how others process information. This provides an alternative, belief-based channel through which decision rights affect behavior.

2.4.2 Delegation and Beliefs About Others

A related literature emphasizes that delegation decisions depend critically on beliefs about others' competence and behavior. Charness et al. (2012) show that individuals' willingness to delegate is sensitive to their beliefs about others' ability and effort. Falk and Kosfeld (2006) demonstrate that control itself affects behavior through expectations and perceived distrust, highlighting the role of beliefs in shaping both actions and outcomes.

This literature implies that the value of transferring decision authority depends on expectations about how others will act. If individuals believe that others will perform worse or interpret information differently, they may prefer to retain control, even when incentives are aligned.

The present paper builds on this insight by focusing specifically on belief updating and information acquisition. I study how individuals form expectations about others' posterior beliefs when priors differ and how these expectations affect their willingness to pay for information. Unlike prior work, which typically focuses on effort or competence, I isolate differences in belief updating as the key dimension of heterogeneity.

The experimental design allows for a direct comparison between individuals' own updating strategies and their expectations of others' strategies. I show that when priors differ, participants anticipate differences in how others interpret signals, which reduces the perceived value of information when decisions are delegated. This mechanism provides a novel explanation for why delegation can affect information demand, even when there is no conflict of interest.

2.5 Cognitive Ability on Belief Updating

The cognitive reflection test (CRT; Primi et al., 2016) is widely used and has been interpreted as a proxy for cognitive ability. It has been found that higher scores in the CRT are related to posterior beliefs closer to Bayes' rule. According to Oechssler et al. (2009), individuals who obtain high scores on the CRT have an average reported belief closer to Bayes' rule than people with lower scores. Similar results were found by Hoppe and Kusterer (2011). Evdokimov and Garfagnini (2023) also mentioned that cognitive ability is associated with how subjects update beliefs; subjects with higher CRT scores demonstrate smaller deviations from the Bayesian benchmark when updating their first-order beliefs.

3 Research Questions

The study explores whether agents consider others' biases in belief updating when priors differ. It also explores if this affects the willingness to share information to persuade other players. More specifically, I ask:

1. How do individuals update their own beliefs and how do they expect others to update when priors differ?

2. How do these beliefs and expectations affect their willingness to pay (WTP) for information when decisions are implemented either by themselves or by another participant?

4 Contributions

This paper is the first experimental study I know that jointly elicits own posteriors, expected posteriors of another agent with possibly different priors, and willingness to pay for information under integrated versus split decision structures in the same environment.

The higher-order updating process was explored experimentally first by Evdokimov and Garfagnini (2022) in the context of sequential learning and common priors. They found that the aggregate distribution of biases is not statistically different (at 5%) between first- and second-order beliefs. My paper explores first, as a marginal contribution, whether these results hold in a different setting; here, the updating was elicited from all possible simultaneous samples with replacement using the strategy method. Second, as a more relevant comparison, I analyze the differences in the higher-order updating under a treatment where participants had different priors. Finally, as the main contribution, I explore the consequences of the hypothesis that participants expect others to update exactly like themselves in the WTP to observe a sample; if this is true, there should not be a difference between guessing the state of the world themselves after observing a sample or letting other participants make the guess.

The delegation of decision-making and the decision to share how much information with others are affected by how people expect others to react to the signals. In this paper, I elicited the WTP to apply the strategies stated as posterior beliefs. If people expect others to use the information in the same way they would, sharing information or acquiring information to make the decision oneself should be the same when the

final decision affects you and others. In a more general setting, if an agent wants to change the beliefs of others, that will depend on how they expect others to use the information observed. For example, if people expect others to be conservative, they will provide more information to change others' beliefs. Also, if the information is costly, they might be discouraged from providing information. It has been found that beliefs can determine the behavior in laboratory experiments (Schotter and Trevino, 2014), and Jin et al. (2021) found that beliefs have an effect on the information provided in a signaling game.

4.1 Methodological Contribution

This study implements the strategy method to evaluate the posterior beliefs of a single person under all possible signals. In the literature, it is common to have sequential belief updating and the direct method (participants observe an actual realization of the process and make a decision). The present setting controls for story effects and is able to compare the decision of each participant in each information condition: prior and signal.

In contrast to Evdokimov and Garfagnini (2022) within-long signal condition, participants are informed that the prior distribution of boxes from another player differs from theirs. Also, the information is presented simultaneously; every signal (sample of marbles) comes from a realization of the state of the world (box selection). The sample of N marbles is common information, and participants report first- and second-order beliefs in the form of strategies.

The strategy method was implemented to explore this question without confounding the effect of different priors as an information source. Thus, participants report their strategies for all possible samples and the expected strategies of someone else. This elicitation method allows us to present the problem without deception and clarifies that other participants' prior probabilities are non-informative. In other studies, the beliefs were elicited using the direct method. However, there is no evidence that applying the strategy method creates a significant difference (Brandts and Charness, 2011).

5 Decision Problem and Optimal Willingness to Pay for Information

This section formalizes the binary decision problem used to elicit willingness to pay for information and derives the corresponding Bayesian benchmark. The environment is characterized by a binary state space,

$$\omega \in \Omega = \{Y, G\},$$

where Y and G denote the two possible states of the world. The agent chooses an action

$$a \in A = \{y, g\},$$

where y represents choosing state Y and g represents choosing state G.

Payoffs depend on whether the selected action matches the realized state. Let

$$u(a, \omega) = \begin{cases} \gamma & \text{if } a = \omega, \\ 0 & \text{if } a \neq \omega, \end{cases} \qquad \text{with } \gamma > 0.$$

Let the prior belief over states be denoted by $\pi_0^\omega = \Pr(\omega)$. Absent any additional information, the optimal action is to select the state with the highest prior probability. The ex ante probability of a correct guess without information is therefore

$$\pi_0^* = \max\{\pi_0^Y, \pi_0^G\},$$

so that the expected payoff based on the prior beliefs is

$$V_0 = \pi_0^* \gamma.$$

Value of the signal Before choosing an action, the agent may acquire a signal σ. The signal structure is described by the conditional probabilities

$$\sigma_\omega^s = \Pr(s \mid \omega).$$

Following the realization of signal s, posterior beliefs are determined by Bayes' rule:

$$\pi_s^\omega = \Pr(\omega \mid s) = \frac{\sigma_\omega^s \pi_0^\omega}{\Pr(s)},$$

where

$$\Pr(s) = \sum_{\omega \in \Omega} \sigma_{\omega}^{s} \pi_{0}^{\omega}.$$

The optimal action after observing s is to choose the state with the highest posterior probability. Accordingly, the conditional probability of making a correct decision after observing s is

$$\pi_{s}^{*} = \max\{\pi_{s}^{Y}, \pi_{s}^{G}\}.$$

The ex ante expected payoff from acquiring the signal is therefore

$$V_{1} = \gamma \sum_{s \in S} \pi_{s}^{*} \Pr(s).$$

Using Bayes' rule, this expression can be written equivalently as

$$V_{1} = \gamma \sum_{s \in S} \max\left\{\sigma_{Y}^{s} \pi_{0}^{Y},\, \sigma_{G}^{s} \pi_{0}^{G}\right\}.$$

This formulation highlights that the value of information is determined by the extent to which the signal increases the probability of choosing the more likely state conditional on each realization.

The optimal willingness to pay for information is given by the instrumental value of the signal, defined as the difference between the expected payoff with and without information:

$$WTP^{*} = V_{1} - V_{0}.$$

Substituting the preceding expressions yields

$$WTP^{*} = \gamma \left(\sum_{s \in S} \max\left\{\sigma_{Y}^{s} \pi_{0}^{Y},\, \sigma_{G}^{s} \pi_{0}^{G}\right\} - \pi_{0}^{*} \right).$$

Calculation with experimental parameters In the decision problem used for the willingness-to-pay elicitation in the experiment, the two states are equally likely: $\pi_{0}^{Y} = \pi_{0}^{G} = 0.5$. And, the payoff from a correct guess is fixed at $\gamma = \$3$. Then, the expected payoff was

$$V_{0} = 0.5 \times 3 = 1.5.$$

In the experimental environment, the state of the world can be a yellow (Y) or green (G) box. The two boxes have five balls and differ in the proportion of yellow

and green balls. Conditional on state Y, the probability of observing a yellow draw is 0.6, whereas under state G the probability is 0.4. The signal is a sample with replacement from a selected box. Hence,

$$\Pr(s \mid Y) = \binom{4}{s}(0.6)^s(0.4)^{4-s},$$

and

$$\Pr(s \mid G) = \binom{4}{s}(0.4)^s(0.6)^{4-s}.$$

The signal consists of a sample of four draws from the selected box, and the realization $s \in S = \{0, 1, 2, 3, 4\}$ denotes the number of yellow draws observed in the sample. With $\pi_0^Y = \pi_0^G = 0.5$,

$$\sum_{s \in S} \max\left\{\sigma_Y^s \pi_0^Y,\, \sigma_G^s \pi_0^G\right\} = 0.648.$$

It follows that

$$V_1 = 3 \times 0.648 = 1.944,$$

and therefore

$$WTP^* = 1.944 - 1.5 = 0.444.$$

Thus, under Bayesian updating and risk neutrality, the optimal willingness to pay for the signal in this decision problem is

$$\boxed{WTP^* = \$0.444.}$$

This benchmark provides the normative reference point against which observed willingness to pay can be compared. Deviations from this value may reflect non-Bayesian updating, misperception of the signal structure, differences in beliefs about how information will be used, or non-instrumental motives such as a preference for retaining decision rights.

5.1 Loss in Willingness to Pay under Integrated and Split Decision Structures

The motivating example highlights that the value of information may depend on whether the individual who acquires information is also responsible for the final decision. In this subsection, I formalize how this distinction affects the willingness to

pay (WTP) for information.

Integrated structure. In the integrated decision structure, the same individual both acquires information and makes the final decision. In terms of the motivating example, this corresponds to Anne (or Ben) deciding whether to hire the consultant and subsequently determining the inventory level for their own product.

Let $\tilde{\pi}^*_{1,s} = \max\{\tilde{\pi}^Y_{1,s}, 1 - \tilde{\pi}^Y_{1,s}\}$ denote the probability that Player 1 selects the correct action after observing signal realization s, based on their (possibly distorted) posterior beliefs. The perceived value of information is given by

$$\tilde{V}_1 = \left(\sum_{s \in S} \tilde{\pi}^*_{1,s} \Pr(s) - \pi^*_0 \right) \gamma.$$

Relative to the Bayesian benchmark WTP^*, any loss in willingness to pay arises from distortions in belief updating. Defining the loss as

$$L_1 = WTP^* - \tilde{V}_1,$$

it follows that

$$L_1 = \gamma \sum_{s \in S} \left(\pi^*_s - \tilde{\pi}^*_{1,s} \right) \Pr(s).$$

Thus, in the integrated structure, the loss in WTP is entirely driven by deviations from Bayesian updating. When beliefs are correct, $\tilde{\pi}^*_{1,s} = \pi^*_s$ for all s, and $L_1 = 0$.

Split structure. In the split decision structure, the individual who decides whether to acquire information is not the one who takes the final action. In the motivating example, Ben decides whether to hire the consultant, while Anne determines the inventory levels.

In this case, Player 1 evaluates the signal based on the expected performance of Player 2. The perceived value of information is

$$\tilde{V}_{2,1} = \left(\sum_{s \in S} \mathbb{E}[\tilde{\pi}^*_{2,s}] \Pr(s) - \pi^*_0 \right) \gamma,$$

where

$$\mathbb{E}[\tilde{\pi}^*_{2,s}] = \tilde{\pi}^*_{2,s} \Pr(a^*_{1,s} = a^*_{2,s}) + (1 - \tilde{\pi}^*_{2,s}) \Pr(a^*_{1,s} \neq a^*_{2,s}).$$

The loss in WTP relative to the Bayesian benchmark is

$$L_{2,1} = WTP^* - \tilde{V}_{2,1} = \gamma \sum_{s \in S} \left(\pi_s^* - \mathbb{E}[\tilde{\pi}_{2,s}^*]\right) \Pr(s).$$

This expression highlights two distinct sources of loss. First, as in the integrated structure, deviations of Player 2's beliefs from the Bayesian benchmark reduce the probability of correct decisions. Second, even if Player 2's beliefs are accurate, Player 1 may anticipate disagreement between their own optimal action and Player 2's action. The term $\Pr(a_{1,s}^* \neq a_{2,s}^*)$ captures this possibility. When disagreement occurs, the effective probability of a correct outcome from Player 1's perspective is reduced, since the final decision does not align with their preferred action.

Comparison. Comparing the two cases, the loss in the split structure satisfies

$$L_{2,1} \geq L_1,$$

whenever Player 1 assigns positive probability to disagreement with Player 2. Intuitively, separating information acquisition from decision-making introduces an additional wedge between the instrumental value of information and its perceived value. Even if information is intrinsically valuable for improving decisions, its willingness to pay is reduced when the decision-maker cannot directly control how that information is used.

In the context of the motivating example, the consultant's forecast is less valuable to Ben when Anne ultimately determines inventory levels, particularly if Ben expects Anne to interpret or use the information differently. This generates a systematic reduction in WTP in the split structure relative to the integrated structure, beyond any distortions arising from non-Bayesian belief updating.

6 Experimental Design

The main situation that the participants have to consider is the following:

Learning Setting

There are two boxes: Yellow (Y) and Green (G). The Yellow box has three yellow marbles and two green marbles, while the Green box has three green marbles

and two yellow marbles. Therefore, the diagnosticity parameter is equal to 3/5 ($P(Y|y) = 1 - P(G|y) = \theta = 0.6$) for a single marble. One box is selected randomly with prior probability $P(Y) \in \pi = \{0.2, 0.5, 0.9\}$, and the signal consists of a random sample with replacement of four marbles (s) drawn from the box. The five possible samples are the number of yellow marbles drawn out of four random samples with replacement: $s \in S = \{0, 1, 2, 3, 4\}$.[5] Participants had to correctly answer six understanding questions after the description of the problem before they could move on to start the experiment.

6.1 Procedures

The experiment was conducted at the University of California, Santa Barbara (UCSB), in the LITE laboratory, in May 2023. I ran twelve sessions with an average size of eleven participants per session. There were 134 participants in total, all students at UCSB: 102 in the DP treatment (ten sessions) and 32 in the CP treatment (two sessions).[6] The average age of the participants was 20.7 years, and 61% were female. The analysis of the WTP in Section 8 also includes four Common Priors sessions with the same design, run in May 2024 (56 participants); they are described there.

The experiment consists of three sections: (1) first- and second-order belief elicitation, (2) signal purchase and strategy implementation, and (3) the CRT and an exit survey. The participants were matched in pairs and belonged to the same pair during the experiment. In Section 1, each pair was in the Common Priors (CP) or Differing Priors (DP) treatment. In the CP treatment, the other participant whose posterior is predicted always shares the participant's prior. In the DP treatment, participants are always asked to predict the posterior of another participant whose prior differs from their own. Treatments were assigned by session: all pairs in a session were in the same treatment. The two CP sessions were run on the same day (May 24, 2023), and the DP sessions on five days between May 18 and May 26. Because the session is the unit of assignment, comparisons between treatments are also tested with randomization inference at the session level, which re-assigns the treatment labels across sessions and asks how often a difference as large as the observed one arises.

[5] There are 2^4 different possible sequences, but the number of marbles of each color is sufficient to calculate the posterior probability.

[6] A pilot session with 8 participants, run before these sessions, used a different version of Sections 2 and 3: participants reported both WTP_i and WTP_{-i}, and the CRT questions were not comparable. It is excluded from all analyses. Including it in the analysis of Section 1 does not change the conclusions.

With two CP sessions out of twelve there are 66 possible assignments, so the smallest attainable p-value is $1/66 \approx 0.015$. Using the strategy method, they had to report their own and another participant's posterior probabilities for each possible sample and prior. In Section 2, they report their willingness to pay to see the signal realized and make a decision according to the strategies stated in Section 1. In each pair, one participant could implement the strategy of the other participant (WTP_{-i}), and the other participant could implement their own (WTP_i). These correspond to the split and integrated decision structures of Section 5, respectively. The combination of these two between-subject treatments in sections one and two generated four possible groups:(CP, WTP_i), (CP, WTP_{-i}), (DP, WTP_i), and (DP,WTP_{-i}). Finally, participants were asked to complete a CRT and an exit survey. The total payment was not revealed until the end of the experiment. The treatments and their conditions are summarized in Table 1. Each section is described in further detail below:

Section	Treatments			
1. First- and Second-Order Belief Elicitation	Common Priors $(P_i(Y), P_{-i}(Y))$		Differing Priors $(P_i(Y), P_{-i}(Y))$	
	$(0.90, 0.90)$		$(0.90, 0.20)$	
	$(0.50, 0.50)$		$(0.50, 0.90)$	
	$(0.20, 0.20)$		$(0.20, 0.50)$	
2. Signal Purchase and Strategy Implementation	Own WTP_i	Another's WTP_{-i}	Own WTP_i	Another's WTP_{-i}
3. CRT and Exit Survey	Demographics and CRT			

Table 1. Experimental design. In Section 1, participants simultaneously report their own and another participant's posteriors for the same realizations. Paired participants are randomly assigned to the CP or DP treatment. The priors they and the other participant have are indicated in parentheses. In Section 2, one participant of the pair reports their WTP_i (integrated structure) and the other their WTP_{-i} (split structure). Finally, all participants complete the CRT and the exit survey.

Section 1: First- and Second-Order Belief Elicitation

Participants' beliefs for each possible sample and prior were elicited using the strategy method. Each participant (i) has to report the probability that the Yellow box was selected given the possible samples ($s \in S$) and stated prior ($P_i(Y) \in \pi$): $\tilde{P}_i(Y|s)$. Also, they have to report the probability stated by another participant $j \neq i$ with prior $P_j(Y)$, denoted $\tilde{P}_{-i}(Y|s)$. The probability of selecting the Yellow box for participant i is $P_i(Y)$, and they know the probability of selecting the Yellow box for participant j:

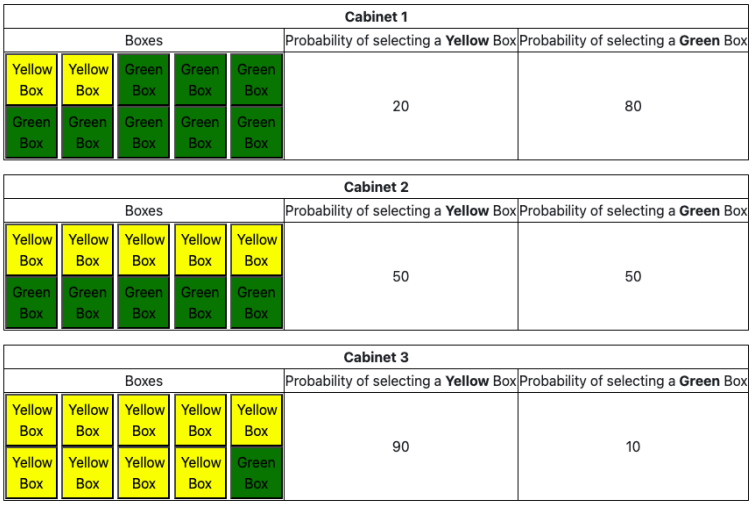


Figure 1. Screenshot of the cabinets used to explain the different priors to the participants.

$P_{-i}(Y)$. In this section there are $(|S| = 5\ \textit{samples}) \times (|\pi| = 3\ \textit{priors}) \times (2\ \textit{players}) =$ 30 observations per participant.

The priors were explained to the participants by presenting three different cabinets (as shown in Figure 1) from where one box was randomly selected for each participant. The boxes came from the same cabinet in the CP treatment, or from different cabinets in the DP treatment. This method allows us to ask for the posterior belief when participants have different priors without deception.

Figure 2 shows an example of the tables with the five samples where the participants reported their strategies. Each prior combination: $(P_i(Y), P_{-i}(Y))$ (see Table 1) was presented in a different table. In this example from the DP treatment, participant one has a prior of 20% and participant two has a prior of 90%. Each participant was asked for their posterior probability of the Yellow box being selected $(\tilde{P}_i(Y|s))$ and the posterior probability of the other participant $\tilde{P}_{-i}(Y|s)$ in the fourth and last columns respectively. Their strategy decisions were made by selecting from a list of 101 numbers between 0 and 100 displayed when clicking on the arrow in each box. They couldn't advance until both participants filled the table.

After stating their strategies for each table, a random process working as described in the learning setting determined a sample s realized for one participant. Therefore they could get paid by their estimation about another participant's strategy or by their own strategy. This practice prevents hedging when participants are reporting their second-order beliefs (Blanco et al., 2010). I use the Binarized Quadratic Scoring

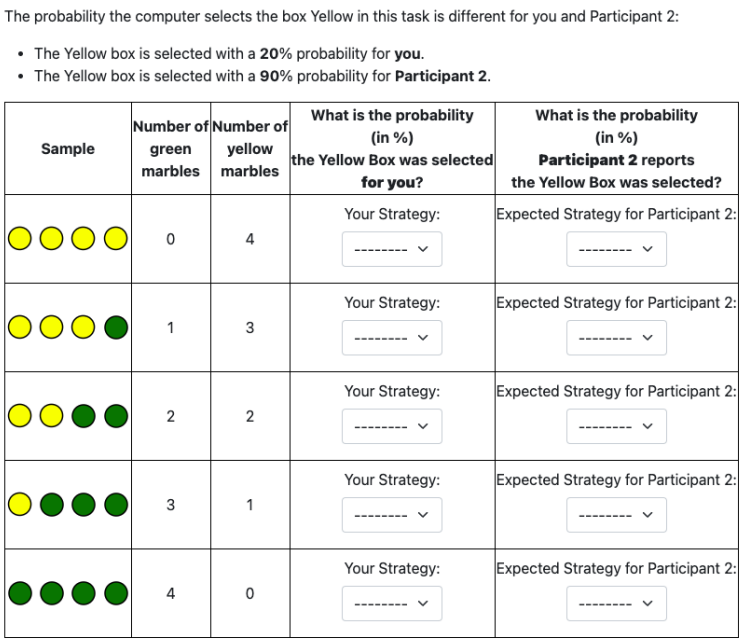


Figure 2. Screenshot of the five possible samples that a participant in the DP treatment with priors $(P_i(Y), P_{-i}(Y)) = (0.20, 0.90)$ has to consider. In the last two columns, they have to indicate $\tilde{P}_i(Y|s)$ and $\tilde{P}_{-i}(Y|s)$, with $s \in \{0, 1, 2, 3, 4\}$ the number of yellow marbles in sample.

Rule (BQSR) (Hossain and Okui, 2013; Erkal et al., 2020) to calculate the payoffs.[7] This mechanism has been used to elicit beliefs in the literature and has the theoretical property that the elicitation is incentive compatible.[8] It is explained to the participants that they must report their true beliefs to maximize their chances of getting up to \$9.[9] Also, the details of how they are incentivized were explained, and some examples presented. The description can be found in Appendix A.

There is a trade-off between complexity, incentives, and accuracy discussed by Charness et al. (2021), and there is evidence that complex descriptions bias the decision towards 50% (Danz et al., 2022).[10] Despite BQSR complexity, this elicitation

[7]Participants make \$3 with probability $1 - (\tilde{P}_i(Y|s) - \mathbb{1}(Y))^2$ from the implementation of their own strategy $(\tilde{P}_i(Y|s))$, and \$3 with probability $1 - (\tilde{P}_{-i}(Y|s) - \tilde{P}_j(Y|s))^2$ from their estimation of the other participant's strategy $(\tilde{P}_{-i}(Y|s))$. $\tilde{P}_j(Y|s)$ refers to the actual strategy of the other participant $j \neq i$.

[8]It has been found that binary lotteries do not induce risk-neutral behavior (Kirchkamp et al., 2021), but the BQSR was more accurate than other options.

[9]They can make \$3 from each of the three tables.

[10]Charness et al. (2021) mentioned that, in some cases, introspection is as good as incentives. However, it is important to maintain the engagement of the participants in the laboratory and incentivize their decisions while keeping the explanation simple.

method allows the implementation of an analogous incentive scheme for first- and second-order beliefs, which makes the strategies stated comparable. Other mechanisms, like probability matching or the interval mechanism, cannot be asked in the same way for their own strategies and those of others. The present implementation of belief elicitation is simple and allows people to read a detailed description of the payoffs with examples.[11] The general description of truthful beliefs maximizing their chances of getting the bonus is correct. This setting maintains a simple description and an incentive-compatible mechanism.

Section 2: Signal Purchase and Strategy Implementation

The BDM mechanism was used to elicit participants' willingness to pay (WTP) to observe the realization s and implement the strategies stated in Section 1. The learning setting is the same as the one described in Section 1, with $P(Y) = 0.5$. Participants are informed that they have the opportunity to purchase a sample of four marbles that is informative of which box was selected. In a second-price auction, participants select a bid b from a list of 31 numbers between 0.00 and 1.50 to bid against the computer that selects a random number $x \sim U[0, 1.5]$. If $b \geq x$, they purchase the information for a price x; if $b < x$, they pay nothing, and no sample is realized.[12] The bid b measures the WTP for the signal, which must equal the difference between the expected value when they receive the information and the expected value of deciding based only on the priors.

The WTP for the signal was elicited in two cases. When participants purchase information and apply their own strategy ($\tilde{P}_i(Y|s)$), b is denoted as WTP_i (integrated structure). When they purchase information but apply the other participant's strategy ($\tilde{P}_j(Y|s)$), b is denoted as WTP_{-i} (split structure). These cases were implemented between subjects; one participant in each pair was asked for their WTP_i, and the other for their WTP_{-i}. Participants were informed that their prediction is based on the strategy reported in Section 1; they were not asked to guess the color of the box after the realized sample s. If the implemented strategy ($\tilde{P}_i(Y|s)$ or $\tilde{P}_j(Y|s)$) is above 0.5, the guess would be the Yellow box; if it is below 0.5, the guess would be the Green box; and a random guess was made if it equals 0.5 or if there was no information purchased through the BDM mechanism. The BDM mechanism was explained to the participants, and they were told that their expected payoffs increase if

[11] The instructions can be found in Appendix A.

[12] The instructions given to the participants can be read in Appendix B.

they report how much extra money they expect to make if the strategy in Section 1 ($\tilde{P}_i(Y|s)$ or $\tilde{P}_j(Y|s)$) is applied to the sample observed.

An extra payoff of \$3 was assigned to each participant if the guess about the box randomly selected was correct. For example, if $x = 0.55 < WTP_{-i} = 0.75$, they purchased the signal. Also, if the Yellow box was selected and the sample $s = 3$ is realized, the guess will be the Yellow box if, for example, $\tilde{P}_j(Y|s) = 0.65$. Then, participant i makes \$3 − \$0.75 = \$2.25 and participant j makes \$3. They make no money for wrong predictions.

The BDM procedure is theoretically robust to risk neutrality.

Section 3: CRT and Exit Survey

Participants could make up to \$2.5 in this section. They were paid \$0.5 for each correct answer in the CRT.

Evdokimov and Garfagnini (2023) found that participants in a sequential updating task underestimate the levels of disagreement regardless of their cognitive ability (measured with a modified CRT) and learning of the state of the world. In their task, participants start with the same prior and learn about the same state of the world sequentially. This is a different setting from the present study, and I expect cognitive ability does not affect how people expect others to update. However, the possibility that cognitive ability and relative self-reported performance could impact higher-order beliefs in the present setting is explored.

7 Results I: Posterior Beliefs (Section 1)

7.1 Posterior Probabilities

Figures 3 and 4 show the average posterior beliefs of the participants in Section 1 of the experiment for the DP and CP treatments, respectively. Each line corresponds to a particular prior and if the posterior is their own strategy or the prediction of another participant's strategy. With a circle shape for each point, the Bayesian benchmark is included in the graph for comparison. A different color (red, green, and blue) is used to differentiate the priors in each posterior belief (20%, 50%, or 90%, respectively). The relative position of the posteriors is predicted according to Grether's model with the estimations from Benjamin (2019) (Figure 7).

Considering a prior of 90% and 50%, participants have posteriors closer to 50% relative to Bayes' rule for all signal levels (the number of yellow marbles in the sample)

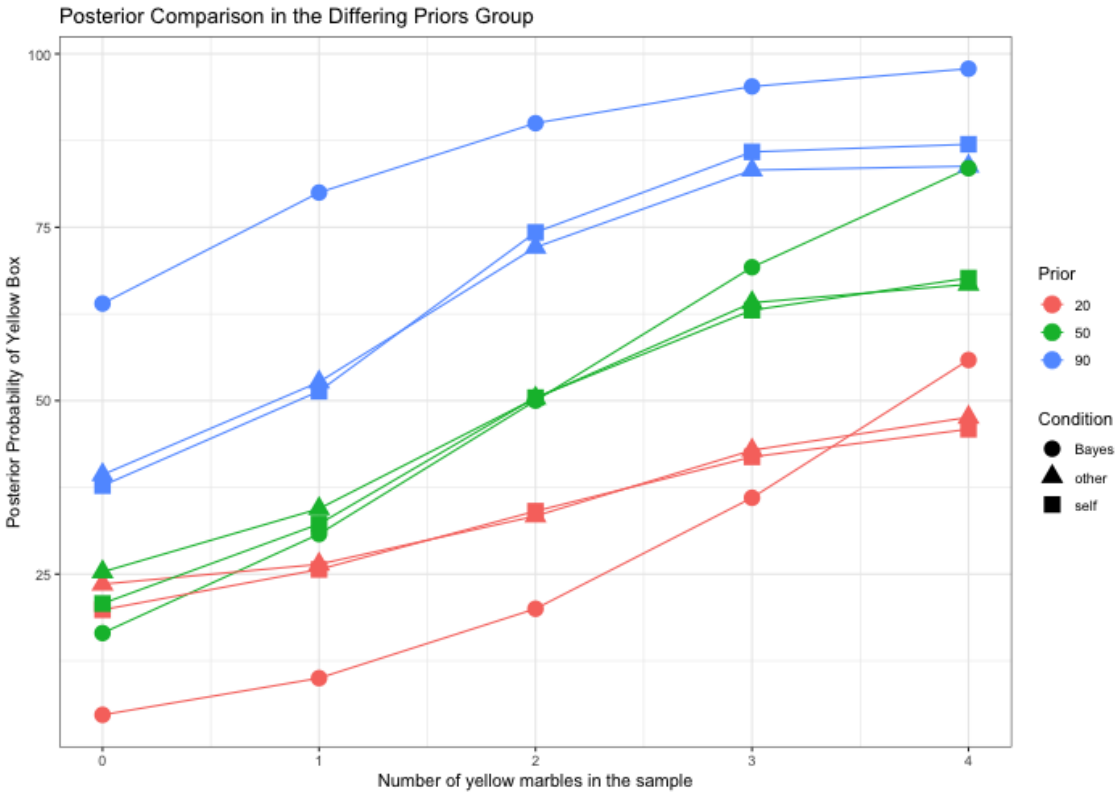


Figure 3. Mean of Posterior Beliefs in the DP Treatment.

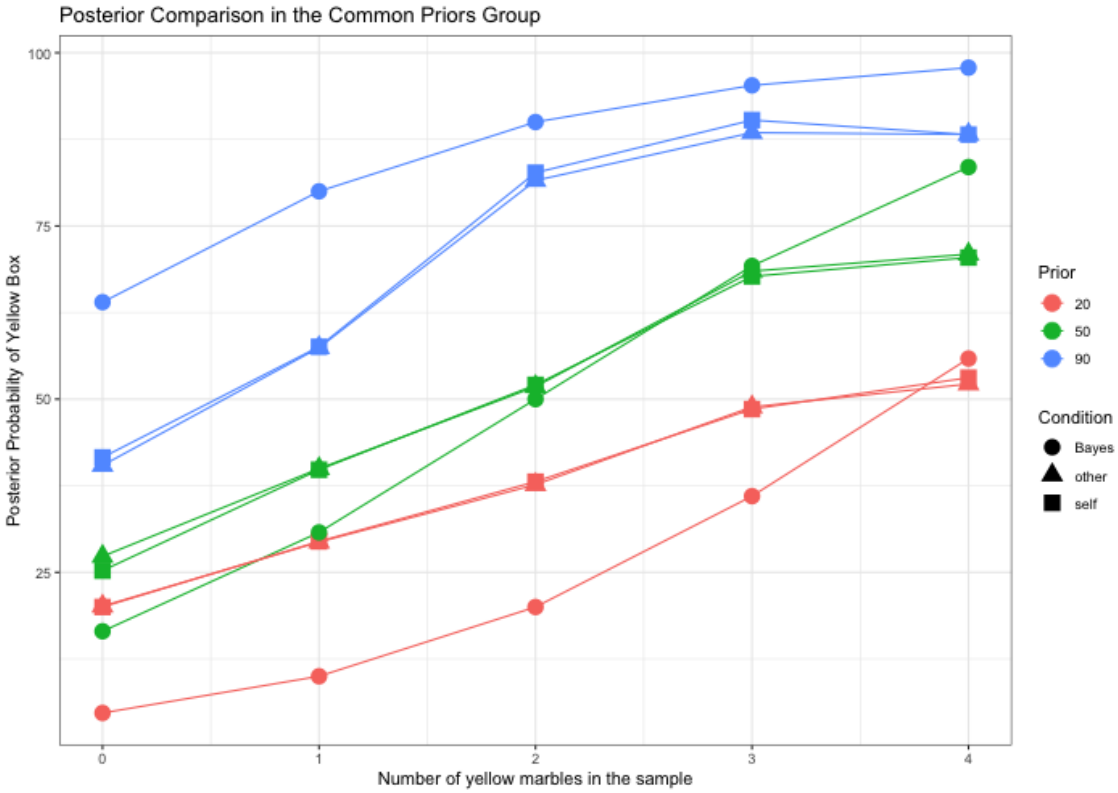


Figure 4. Mean of Posterior Beliefs in the CP Treatment.

except for the non-informative signal of two yellow marbles in the sample. The same patterns follow with the 20% prior except for the sample with four yellow marbles. In this case, participants are over-inferring on average. This is consistent with the model's predictions based on estimations from previous studies.

The squares and triangles denote the posterior beliefs when they report their own beliefs or the beliefs of another participant. There is no important difference between these two conditions, which means that people, on average, estimate that others will behave like themselves.

When participants have to predict the behavior of another participant with the same prior, they predict others to be like themselves. However, when they have different priors, they predict, on average, the same behavior but with a larger variance. This can be observed in the box plot of the differences between first- and second-order posteriors. Figures 5 and 6 show these differences for the DP and CP treatments, respectively. In the DP treatment, it seems that distances do not depend on the information conditions. Except for the non-informative signal of two yellow marbles and the uniform prior of 50%, where most people predicted a posterior of 50% for the

other participant, which was their strategy as well. In the CP treatment, the vast majority of differences were zero, which means that people expect others to behave exactly like themselves. This exploratory analysis with box plots is useful, but a more specific regression to test this hypothesis was run as well.

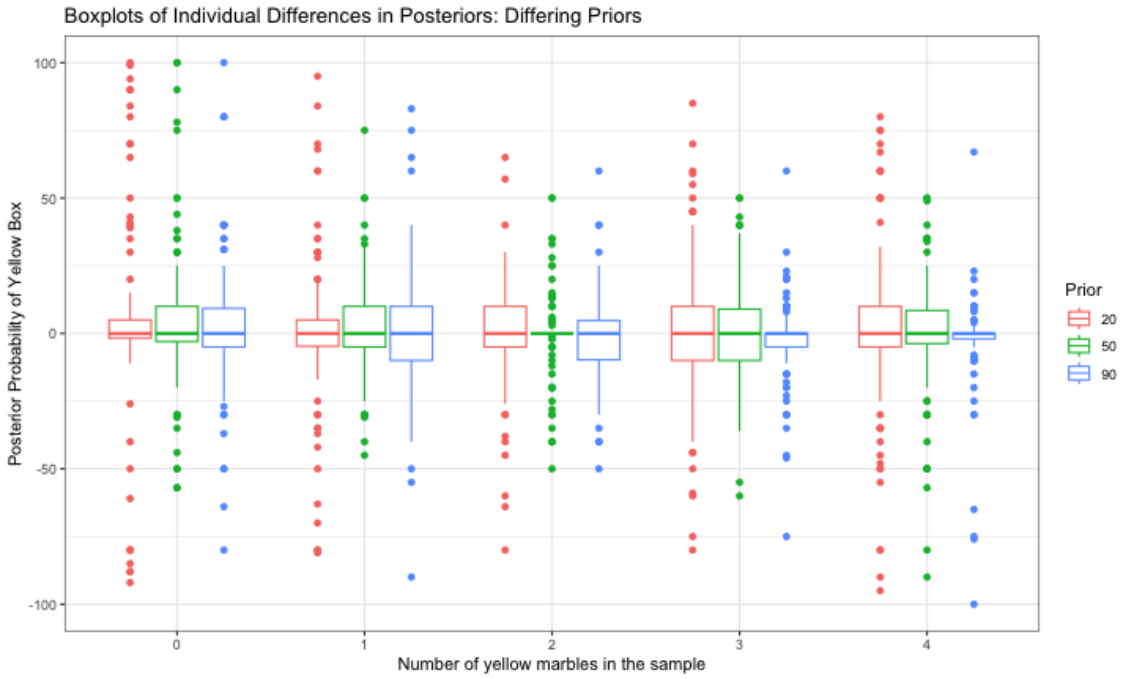


Figure 5. Boxplots of the differences between the posterior probability reported and predicted for another participant in the DP treatment. These participants didn't face the same prior at the same time, but since they were asked about their strategies and others in different combinations of priors, the first- and second-order beliefs can be compared directly.

The regression in Table 2 shows no systematic effect on the differences between $\tilde{P}_i(Y|s)$ and $\tilde{P}_{-i}(Y|s)$. For this regression, the answers were clustered at the participant level.

There is no effect on the posterior belief that this was for themselves or an expectation. Posteriors are 4.1 percentage points higher in the CP treatment ($p = 0.048$ with standard errors clustered by participant), but this difference is not significant when the treatments are re-assigned across sessions ($p = 0.14$), so there is no robust evidence of a difference in the level of posteriors between treatments. This regression doesn't test if there was a difference between first- and second-order posterior beliefs, but only if that difference is constant across information conditions. It is still possible that the differences in the first- and second-order beliefs follow a pattern for different

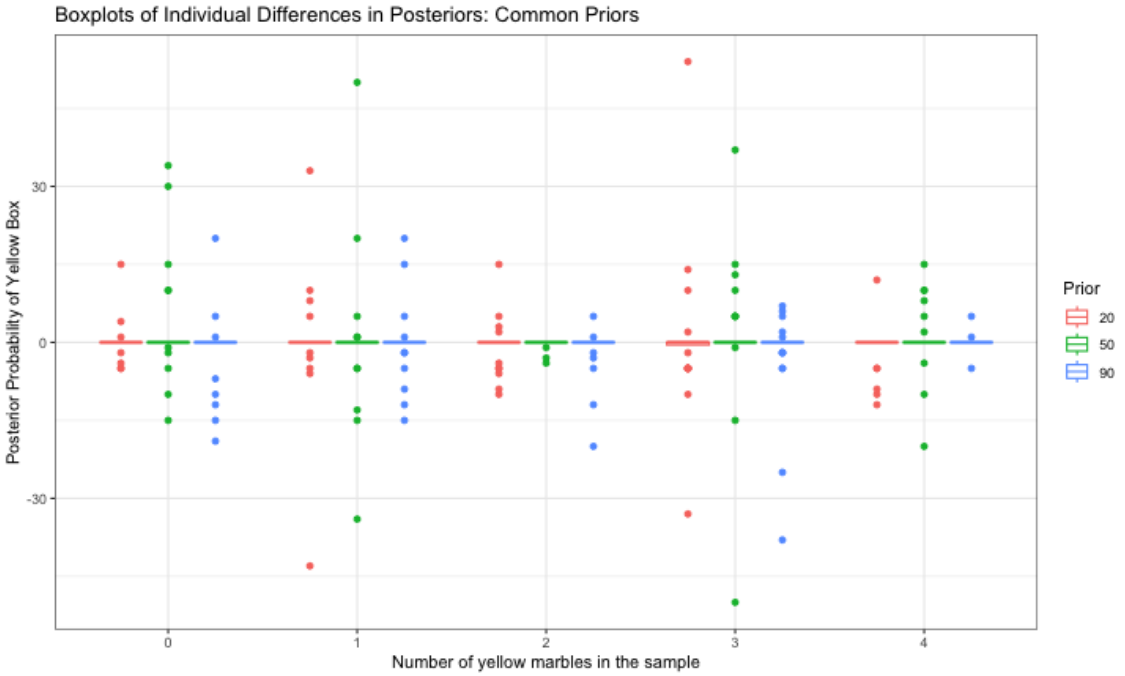


Figure 6. Boxplots of the differences between the posterior probability reported and predicted for another participant in the CP treatment. They were asked for their posterior belief and that of another participant with exactly the same prior simultaneously.

priors and signals. The estimated coefficients of model (2) are compared in the next section to analyze this.

The regression in Table 3 takes the absolute difference as the dependent variable. For this regression, the answers were clustered at the participant level. The absolute difference $|\tilde{P}_i(Y|s) - \tilde{P}_{-i}(Y|s)|$ is 11.2 percentage points smaller in the CP treatment. This difference is also significant at the session level ($p = 0.015$, the smallest attainable with two CP sessions).

Table 2. Regression of the Reported Posteriors on the Type of Report and Treatment.

	Dependent variable: posterior (%)
	(1)
Own posterior	-0.387
	(0.592)
CP treatment	4.135**
	(2.086)
Prior (%)	0.471***
	(0.030)
No. yellow marbles	10.507***
	(0.684)
Constant	3.500**
	(1.754)
Observations	4,020
R^2	0.388

Note: *p<0.1; **p<0.05; ***p<0.01
Standard errors clustered by participant. Own posterior = 1 for the participant's own posterior and 0 for the predicted posterior of the other participant. Treatments were assigned by session; re-assigning them across the 12 sessions (randomization inference) gives $p = 0.136$ for the CP coefficient.

7.2 Grether's Model

To analyze biases in higher-order belief updating, I use the logistic regression model used in the original Grether (1980) paper. I will change the notation of the model to capture the deviations from Bayes' rule between first- and second-order beliefs. The reported probability by i about the expected posterior probability of $j \neq i$ is denoted by $\tilde{P}_{-i}(Y|s)$. The model for belief updating is:

Table 3. Regression of the Absolute Difference between Own and Predicted-Other Posteriors.

	Dependent variable: $\lvert\tilde{P}_i(Y\lvert s) - \tilde{P}_{-i}(Y\lvert s)\rvert$
	(1)
CP treatment	-11.188***
	(1.270)
Prior (%)	-0.072***
	(0.020)
No. yellow marbles	-0.566**
	(0.285)
Constant	18.814***
	(1.989)
Observations	2,010
R^2	0.083

Note: *p<0.1; **p<0.05; ***p<0.01

Standard errors clustered by participant. Re-assigning treatments across the 12 sessions (randomization inference) gives $p = 0.015$ for the CP coefficient, the smallest p-value attainable with 66 possible assignments.

$$\tilde{P}_i(Y|s) = \frac{\left(P_i(s|Y)\right)^{\beta_i^l}\left(P_i(Y)\right)^{\beta_i^p}}{\left(P_i(s|Y)\right)^{\beta_i^l}\left(P_i(Y)\right)^{\beta_i^p} + \left(P_i(s|G)\right)^{\beta_i^l}\left(P_i(G)\right)^{\beta_i^p}} \tag{1}$$

However, to abstract from the denominator and easily estimate the parameters in a regression, the logarithmic transformation of the posteriors is considered:

$$\frac{\tilde{P}_i(Y|s)}{\tilde{P}_i(G|s)} = \left(\frac{P_i(s|Y)}{P_i(s|G)}\right)^{\beta_i^l}\left(\frac{P_i(Y)}{P_i(G)}\right)^{\beta_i^p}$$

$$\log\left(\frac{\tilde{P}_i(Y|s)}{\tilde{P}_i(G|s)}\right) = \beta_i^l \log\left(\frac{P_i(s|Y)}{P_i(s|G)}\right) + \beta_i^p \log\left(\frac{P_i(Y)}{P_i(G)}\right)$$

The expected probability of participant i with respect to participant j is $\tilde{P}_{-i}$. When referring to first-order beliefs of i the notation will be $\tilde{P}_i(B|s)$, $P_i(B)$ and $P_i(s|B)$, $B \in \{Y, G\}$. The model to adjust to the data is, therefore:

$$\log\left(\frac{\tilde{P}_i(Y|s)}{\tilde{P}_i(G|s)}\right) = \beta_i^l \log\left(\frac{P(s|Y)}{P(s|G)}\right) + \beta_i^p \log\left(\frac{P_i(Y)}{P_i(G)}\right) \tag{2}$$

The parameter β^p_{-i} measures the weight given to the prior probability assigned to $j \neq i$ when i has to report $\tilde{P}_{-i}(Y|s)$, and β^l_{-i} the weight given to the likelihood. Benjamin (2019) found that the parameters β^p_i and β^l_i (first-order beliefs) are affected by the sample size and if the information was presented simultaneously or sequentially. The present methodology implements a simultaneous and constant (N=4) sample. The estimation of both parameters will happen simultaneously.[13]

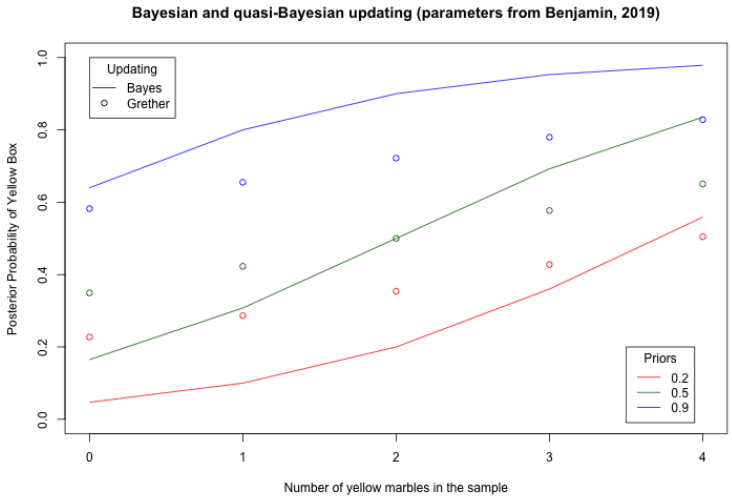


Figure 7. Posterior distribution as a function of the signal and the priors according to Grether's quasi-Bayesian model (Equation (1)). The parameters $\beta^l_i = 0.383$ and $\beta^p_i = 0.434$ were taken from the regressions in Benjamin (2019) for the simultaneous and incentivized experiments. These parameters, and the likelihoods and priors used in the experiments, were used to plot the dots.

To answer the research question, it is necessary to test if $\beta^p_{-i} = \beta^p_i$ and $\beta^l_{-i} = \beta^l_i$. This corresponds with the parameters estimated in the CP and DP treatments. Figure 7 shows the estimated posterior ratios according to the model in Equation (2) and the parameters used in the present methodology ($\theta = 0.60$, $N = 4$). In the figure, the expected posterior for a Bayesian agent and the adjusted parameters in Benjamin (2019) for simultaneous and incentivized experiments ($\beta^l_i = 0.383$, $\beta^p_i = 0.434$) are shown.

[13] Benjamin (2019) first estimated the parameter β^l as a function of sample size and the timing of the information and used those estimations to estimate β^p so there was no bias since β^l is not constant. Here, the sample has a constant size and is simultaneous, and (β^l, β^p) will be estimated at the same time.

7.2.1 Grether's Model Estimations

The regressions in Table 4 evaluate if there is a systematic difference between the parameters in the quasi-Bayesian model (Equation (2)) for the CP and DP treatments and if it was their strategy $\tilde{P}_i(Y|s)$ or the expected strategy for another participant $\tilde{P}_{-i}(Y|s)$. All the parameters deviate from Bayes' rule: all of them are statistically different from 1. Participants exhibit conservatism and base-rate neglect. It is also apparent that the differences between the parameters for $\tilde{P}_i(Y|s)$ and $\tilde{P}_{-i}(Y|s)$ are not important.

The values of the parameters estimated in this experiment are consistent with the results from Benjamin (2019). In this case, it was not necessary to estimate β^l first since the sample size and information provision were constant; Benjamin (2019) found that these conditions affect the value of the parameters. Both parameters were estimated simultaneously. Also, the condition with prior 50% was evaluated for each participant, which allows for a better estimation of β^l.

Table 4. Grether's Model Regressions by Prior Treatment and Report.

	Dependent variable: log posterior odds			
	Other CP	Other DP	Own CP	Own DP
	(1)	(2)	(3)	(4)
β^p	0.672***	0.537***	0.697***	0.619***
	(0.100)	(0.052)	(0.104)	(0.054)
β^l	0.698***	0.541***	0.674***	0.621***
	(0.134)	(0.067)	(0.133)	(0.066)
Constant	-1.144***	-1.145***	-1.114***	-1.363***
	(0.307)	(0.133)	(0.314)	(0.131)
Observations	480	1,530	480	1,530
R^2	0.253	0.189	0.249	0.246

Note: *p<0.1; **p<0.05; ***p<0.01
Standard errors clustered by participant. "Own" columns use participants' own posteriors; "Other" columns use their predictions of the other participant's posteriors.

7.2.2 Individual Analysis

The previous aggregate analysis calculates the parameters as if they were the same for all participants. This is useful to have an idea of the average behavior. However, there could be heterogeneity in the biases that participants have. The regression in

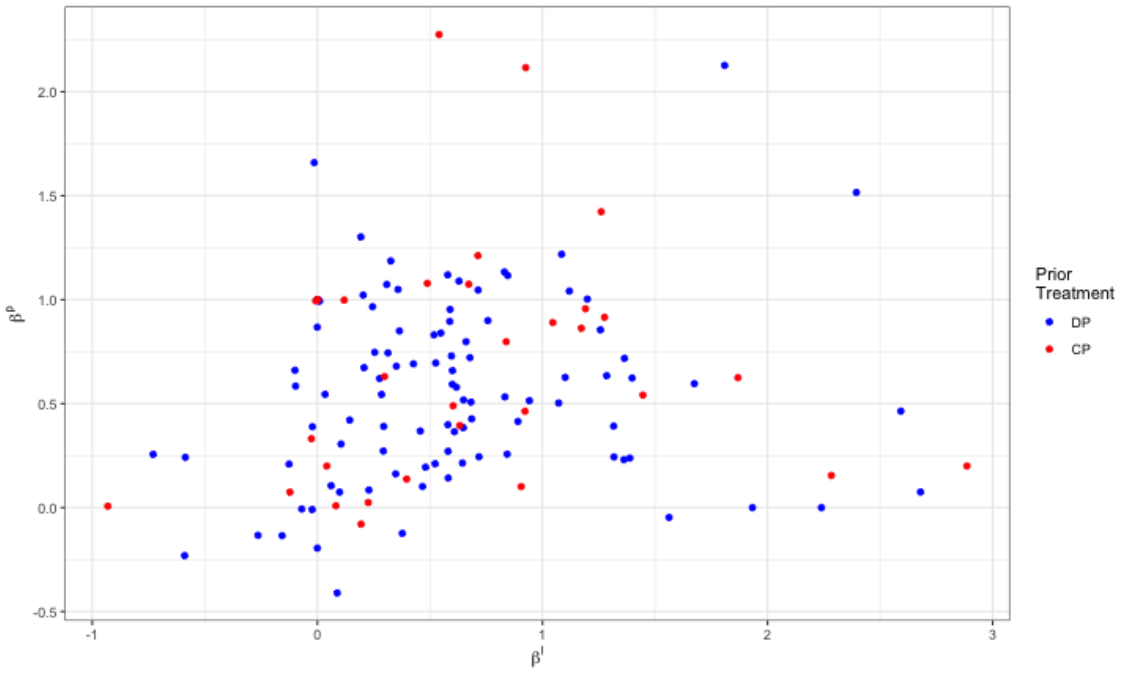


Figure 8. Grether's model estimated at the individual level. Treatments: Common Priors (CP) and Differing Priors (DP).

Equation (2) was run for each participant to know the biases they have. Figure 8 plots the coefficients for each participant. The correlation between the coefficients was 0.11 and not statistically significant ($p = 0.22$). Therefore, we cannot classify people in one dimension with different degrees of Bayesianism.[14] If the correlation were positive, it could be possible to classify people by their level of Bayesianism: they have high or low β's. If the correlation were negative, it could be possible to classify people as conservative or BRN. However, the heterogeneity of the biases seems uniformly distributed in the unit square. The parameters of participants in the CP and DP treatments are also differentiated in Figure 8 with no clear distinction between them.

Table 5. Correlation of the Individual Grether Coefficients (DP Treatment).

	β_i^p	β_{-i}^p	β_i^l	β_{-i}^l
β_i^p	1	0.350^{***}	0.158	0.157
β_{-i}^p	0.350^{***}	1	0.004	0.001
β_i^l	0.158	0.004	1	0.804^{***}
β_{-i}^l	0.157	0.001	0.804^{***}	1

Pearson correlations, 102 participants in the DP treatment. *p<0.1; **p<0.05; ***p<0.01.

[14] This is relevant when analyzing other models, like that of Epstein et al. (2008), where the biases are a one-dimensional deviation from Bayes' rule.

The coefficients of Grether's model were also calculated conditioned on the strategy being their own or another participant's (β_i and β_{-i}). These correlations only include the DP treatment since in the CP treatment, the expectation of others answering in the same way is very strong, as seen in Figure 6. Table 5 reports the correlations.[15]

Consistent with the previous analysis, the correlations were positive and significant between β_i^p and β_{-i}^p (0.35), and between β_i^l and β_{-i}^l (0.80). Projection is therefore much stronger for the weight on the signal than for the weight on the prior. Also, the correlation was not significant between β^p and β^l regardless of whether they are for first- or second-order beliefs. These correlations were close to zero (at most 0.16 in absolute value, for β_i^p with β_i^l and with β_{-i}^l) and none was significantly different from zero.

The correlations show a positive relationship between the biases that participants displayed and what they predicted for others. They also show a low correlation between BRN and conservatism, which means that there is no evidence of people having a stronger BRN bias and also having a larger (or smaller) level of conservatism. The correlations inform us about the linear relationship, but analyzing the differences between the parameters can tell us if people expect others to be more biased.

Participants in this experiment might exhibit a bias blind spot (Pronin et al., 2002). This means that the biases calculated for others (β_{-i}) is larger than for themselves ($|1 - \beta_i| < |1 - \beta_{-i}|$). The hypothesis $\beta_i > \beta_{-i}$ can be considered since participants were generally downwards biased ($\beta_i < 1$). No significant difference was found between β_i^l and β_{-i}^l, nor between β_i^p and β_{-i}^p in the CP treatment, nor between β_i^p and β_{-i}^p in the DP treatment. However, a marginally significant difference of 0.080 between β_i^l and β_{-i}^l was found with a paired t-test in the DP treatment (one-sided $p = 0.029$; two-sided $p = 0.058$). The mean for β_i^l is larger than β_{-i}^l, consistent with a larger bias expected from others in the weight given to the likelihood of the signals.

The previous analysis of the individual parameters suggests that, in the DP treatment, people expect others to be slightly more prone to conservatism than themselves. Also, there is no relationship between BRN and conservatism.

[15] When the CP treatment is included, the correlations between first- and second-order beliefs are larger, and no significant correlation is found among the other coefficients. These results are included in Appendix C.

7.2.3 Treatment Effects on β

Figures 9 and 10 show the cumulative distributions of β^l and β^p by prior treatment and by whether the parameter describes the participant's own strategy or the predicted strategy of the other participant. Figure 11 shows the distribution of the individual differences $\beta^l_{-i} - \beta^l_i$ by prior treatment.

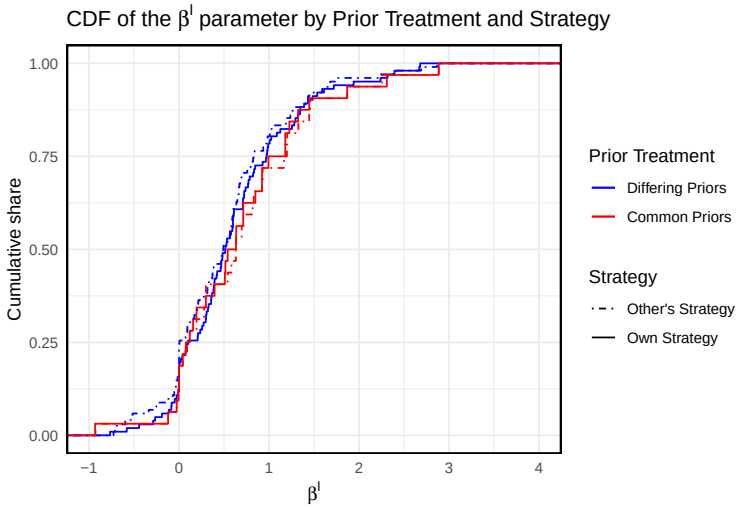


Figure 9. Distribution of β^l by prior treatment and strategy (own or other's).

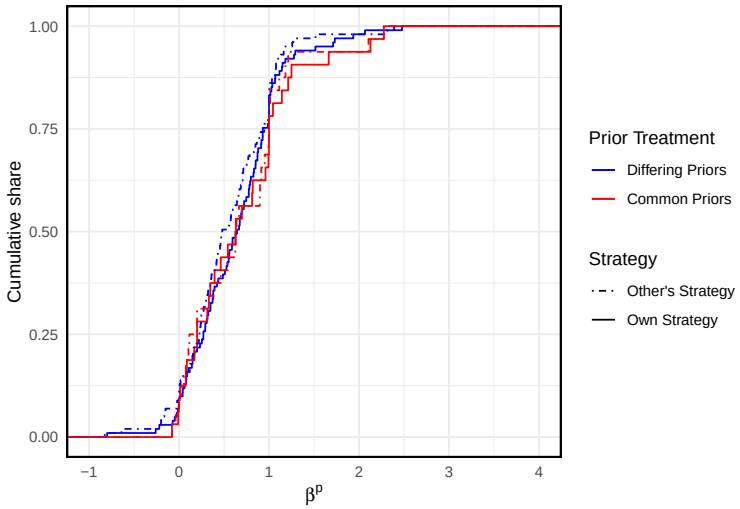


Figure 10. Distribution of β^p by prior treatment and strategy (own or other's).

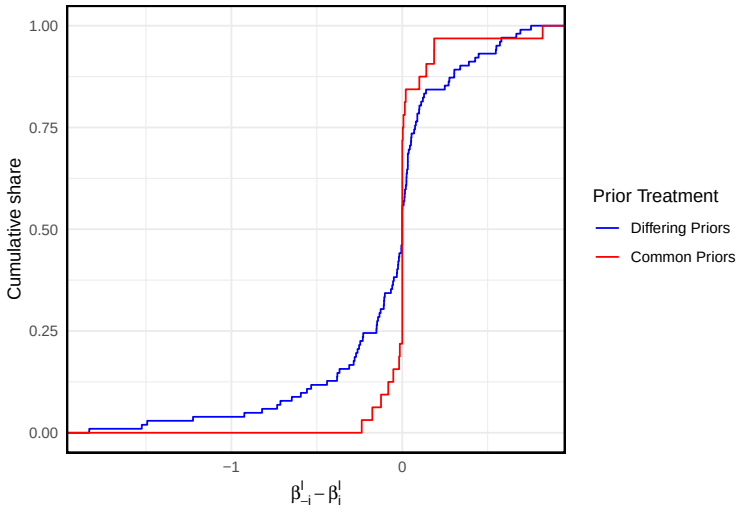


Figure 11. Differences β^l by prior treatment (CP vs. DP). A negative number means that the participant expected a smaller weight of the signal for the other participants relative to themselves. β^l_{-i} represents expected bias from others, β^l_i self-bias.

7.3 Robustness: Between-Subject Elicitation and the Monotonicity Check

Two additional sets of sessions, run after the main experiment, examine whether these belief results depend on how posteriors were elicited. They are not part of the main sample, and Appendix D reports the comparisons in detail.

Between-subject elicitation. In the main experiment, participants reported their own posterior and their prediction of the other participant's posterior side by side, which could make it easy to copy their own answer. In four sessions (52 participants, November–December 2023), all pairs shared a common prior, and each participant reported either their own posteriors or their prediction of the partner's posteriors, but not both. Predictions of the partner's posteriors were indistinguishable from the posteriors that other participants reported for themselves (difference of -1.5 percentage points, $p = 0.54$), and the Grether parameters of the two groups did not differ ($p = 0.81$ for β^l and $p = 0.60$ for β^p). Posteriors and parameters were also similar to those in the CP treatment of the main sample (Appendix Table D1). Participants therefore expect others to update like themselves even when they cannot copy their own answer.

Monotonicity check. In four Common Priors sessions (56 participants, May 2024), the design was the same as in the CP treatment, except that the interface flagged

strategy tables that were not monotone in the number of yellow marbles; 22 participants were flagged at least once. All tables became monotone (compared with 72% in the CP treatment of the main sample), and the weight on the signal moved closer to the Bayesian value of 1 (median β_i^l of 0.95 compared with 0.59; $p = 0.004$). Participants also copied their own posteriors less often when predicting the partner's (58% of identical reports compared with 74%; $p < 0.001$). The average distance between posteriors and Bayes' rule fell from 20.1 to 17.0 percentage points, but this difference is not statistically significant ($p = 0.13$; Appendix Table D3).

8 Results II: Willingness to Pay (Section 2)

8.1 Sample

The analysis of the WTP pools the main sample with the four Common Priors sessions run in May 2024. These sessions used the same Section 2 as the CP treatment: one participant in each pair bid to implement their own strategy and the other bid to implement the partner's strategy, at the 50% prior; participants also answered an additional question about their average bet. The monotonicity check changed how participants updated (Section 7.3), but it did not change their bids. Bids for one's own strategy were \$0.67 in 2024 and \$0.64 in 2023 ($p = 0.84$), bids for the partner's strategy were \$0.64 and \$0.58 ($p = 0.72$), and the gap between them did not differ between years ($p = 0.90$; Appendix Table D3). Pooling raises the number of CP pairs from 16 to 44, so the pooled WTP sample has 190 participants: 102 in the DP treatment and 88 in the CP treatment.

The between-subject sessions cannot be pooled. All their pairs shared a common prior, so they contain no variation between the DP and CP treatments, and their bids were elicited at three different priors by participants whose role in Section 2 was tied to the type of posterior they reported in Section 1. They are used separately in Section 8.4.

Because the monotonicity check shifted participants' Grether parameters without changing their bids, the analyses that relate WTP to individual updating (Section 8.3) use only the main sample.

8.2 WTP by Prior Treatment and Decision Structure

Figure 12 shows the average WTP by prior treatment and decision structure. All four averages are above the Bayesian benchmark of \$0.444, and 74% of individual bids

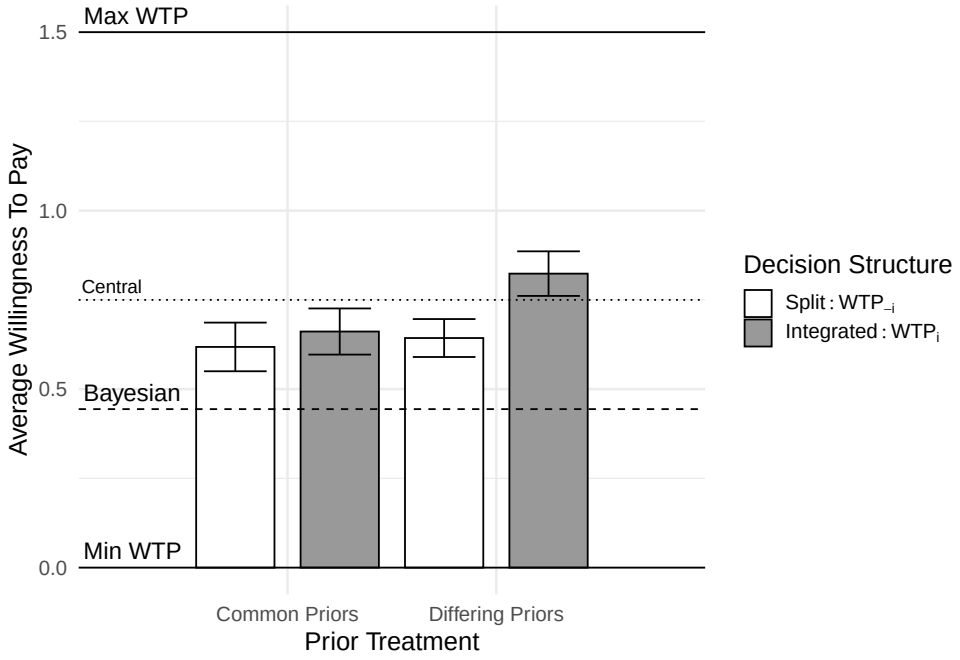


Figure 12. Average WTP by prior treatment and decision structure (WTP_i: integrated; WTP_{-i}: split), pooled WTP sample (main sample and the 2024 Common Priors sessions; 190 participants). Error bars show standard errors. The dashed line marks the Bayesian benchmark $WTP^* = \$0.444$; the solid lines mark the minimum and maximum possible bids, and the dotted line the central bid ($0.75).

exceed it, so participants value the signal more than a risk-neutral Bayesian agent would. Bids are highest when participants in the DP treatment implement their own strategy (\$0.82, SE 0.06), compared with \$0.64 (SE 0.05) when they implement the other participant's strategy (Welch t-test, $p = 0.03$; Mann-Whitney $p = 0.04$). In the CP treatment, bids under the two decision structures are close (\$0.66 and \$0.62; $p = 0.65$). Comparing across treatments, bids for one's own strategy are higher in the DP treatment than in the CP treatment (\$0.82 and \$0.66; $p = 0.08$; $p = 0.10$ when the treatments are re-assigned across the 16 sessions), while bids for the partner's strategy are similar (\$0.64 and \$0.62; $p = 0.77$; session level $p = 0.81$). The same pattern appears within pairs, since each pair had one participant in each decision structure: the own-minus-partner difference averages \$0.18 in DP pairs ($p = 0.01$; 51 pairs) and \$0.04 in CP pairs ($p = 0.62$; 44 pairs).

Whether the premium is larger under differing priors depends on the difference between the two gaps. Column (1) of Table 6 estimates it as the interaction between implementing one's own strategy and the DP treatment: 0.14 (robust SE 0.13; $p = 0.27$; 95% confidence interval $[-0.11, 0.38]$). Comparing the pair differences directly gives the same estimate ($p = 0.21$; one-sided $p = 0.11$). The interaction does not change when the 2024 sessions are allowed to have their own gap (column 2; the gap in those sessions does not differ from the 2023 CP sessions, $p = 0.90$) or when CRT controls are added (column 3: 0.16; $p = 0.20$). Pooling makes the test more precise than with the main sample alone (column 4: 0.12; $p = 0.50$): the smallest difference between the gaps that the design can detect with 80% power falls from about \$0.50 to \$0.35. It is still larger than the observed difference. Randomization inference at the session level, which accounts for the assignment of treatments by session, gives the same conclusion: $p = 0.29$ for the interaction in column (1) and $p = 0.50$ in column (4).

As a robustness check, the DP sessions can also be compared with every session in which priors were common, including the between-subject sessions, using their bids at the 50% prior (Section 8.4). Bids for one's own strategy are \$0.82 in the DP sessions and between \$0.61 and \$0.67 in each of the other three groups of sessions (session level $p = 0.05$; 20 sessions), while bids for the partner's strategy are between \$0.58 and \$0.64 in all groups ($p = 0.70$). The difference between the two gaps remains insignificant ($p = 0.20$). Because the between-subject sessions differ in design and date, this comparison complements but does not replace the main test.

The evidence is therefore consistent with exposure to differing priors raising the

value of implementing one's own strategy, since the premium appears only in the DP treatment and comes entirely from bids for one's own strategy, but the difference between treatments is not statistically established. Because treatments were assigned by session, it should be read as suggestive rather than causal evidence.

It is important to remember that participants are not making the decision again; the strategies they reported in Section 1 are implemented mechanically. The premium for implementing one's own strategy in the DP treatment therefore does not come from the opportunity to revise the decision.

One of the hypotheses was that people who expect others to update like them would have $WTP_i \approx WTP_{-i}$. This holds in the CP treatment, where participants also expect others to update exactly like themselves (Figure 6). In the DP treatment, the expectation about others is on average the same but more dispersed, and bids for one's own strategy are higher.

8.3 WTP and Individual Updating

This subsection uses the main sample only (134 participants). In the 2024 sessions, the monotonicity check raised the weight on the signal without changing bids, so pooling them would mix two ways of measuring updating. For example, when the regression of own-strategy bids in Table 8 is estimated on the pooled sample, the DP coefficient becomes 0.21 ($p = 0.01$), but it falls to 0.19 ($p = 0.14$) once an indicator for the 2024 sessions is added. In the main sample, the DP coefficient is 0.18 ($p = 0.17$; session level $p = 0.33$). Its size barely changes when the perceived value of the signal is controlled for, so the higher own-strategy bids in the DP treatment are not explained by differences in how participants update.

Table 7 relates bids to the individual Grether parameters. Among participants who implement their own strategy, WTP increases with the weight on the signal, β_i^l (0.19, $p = 0.005$), and decreases with the weight on the prior, β_i^p (-0.16, $p = 0.06$). A larger β^l means that participants are less conservative, so new information moves their posteriors more; a larger β^p means that they rely more on the prior (i.e., they exhibit less BRN), so new information matters relatively less. Both signs are consistent with participants valuing the signal according to how much it moves their own beliefs. Bids to implement the other participant's strategy are not related to the parameters of the participant's predictions of the other (column 2); they are related only to the participant's own weight on the prior (column 4: -0.16, $p = 0.04$). Because β_i and β_{-i} are highly correlated (Table 5), the cross-specifications in columns (3) and (4)

Table 6. Willingness to Pay by Prior Treatment and Decision Structure.

	Dependent variable: WTP ($)			
	Pooled	Pooled	Pooled	Main sample
	(1)	(2)	(3)	(4)
Own strategy	0.043	0.059	0.021	0.059
	(0.094)	(0.158)	(0.094)	(0.158)
DP treatment	0.025	0.059	0.004	0.059
	(0.086)	(0.129)	(0.087)	(0.129)
Own strategy × DP	0.137	0.121	0.162	0.121
	(0.125)	(0.179)	(0.125)	(0.178)
2024 session		0.053		
		(0.144)		
Own strategy × 2024 session		-0.025		
		(0.197)		
CRT better than partner			0.032	
			(0.068)	
CRT score			-0.047	
			(0.029)	
Constant	0.618***	0.584***	0.700***	0.584***
	(0.068)	(0.118)	(0.088)	(0.117)
Observations	190	190	190	134
R^2	0.037	0.038	0.052	0.049

Note: *p<0.1; **p<0.05; ***p<0.01

Heteroskedasticity-robust (HC1) standard errors. Own strategy = 1 for participants who bid to implement their own strategy (WTP_i, integrated structure) and 0 for those who bid to implement the other participant's strategy (WTP_{-i}, split structure). Pooled: main sample plus the 2024 Common Priors sessions; 2024 session = 1 for these sessions. Within-treatment Welch t-tests of own vs. other's WTP (pooled sample): DP gap = 0.180 (p = 0.030); CP gap = 0.043 (p = 0.647). Treatments were assigned by session; re-assigning them across sessions (randomization inference) gives p = 0.286 for Own strategy × DP in column (1) (16 sessions) and p = 0.500 in column (4) (12 sessions).

should be read with caution.

Table 8 summarizes the parameters in a single measure: the perceived value of the signal, $\tilde{V}$, implied by the Grether parameters of the strategy that is implemented (Section 5). $\tilde{V}$ predicts WTP for one's own strategy (0.53, $p = 0.03$) but not for the other participant's strategy (0.09, $p = 0.70$). Sensitivity to information therefore predicts WTP when participants implement their own strategy, as in the integrated structure of Section 5, while bids for someone else's strategy appear unrelated to the beliefs participants hold about how that person updates.

Table 7. Willingness to Pay and Individual Grether Parameters (Main Sample).

	Dependent variable: WTP ($)			
	Own strategy	Other's strategy	Own strategy	Other's strategy
	(1)	(2)	(3)	(4)
β_i^p	-0.160*			-0.159**
	(0.084)			(0.075)
β_i^l	0.186***			0.019
	(0.064)			(0.068)
β_{-i}^p		-0.095	-0.250**	
		(0.077)	(0.105)	
β_{-i}^l		-0.014	0.124*	
		(0.067)	(0.065)	
Constant	0.736***	0.693***	0.835***	0.727***
	(0.081)	(0.067)	(0.086)	(0.067)
Observations	67	67	67	67
R^2	0.126	0.022	0.118	0.050

Note: *p<0.1; **p<0.05; ***p<0.01
Heteroskedasticity-robust (HC1) standard errors. Own strategy: participants who bid to implement their own strategy (WTP_i, integrated structure); Other's strategy: participants who bid to implement the other participant's strategy (WTP_{-i}, split structure). β_i: parameters of the participant's own posteriors; β_{-i}: parameters of their predictions of the other participant.

8.4 WTP and the Prior

In the main experiment, the WTP was elicited only at the 50% prior. In the between-subject sessions, each participant bid once at each of the three priors (20%, 50%, and 90%), which shows whether bids respond to the value of the signal. The Bayesian

Table 8. Willingness to Pay, the Perceived Value of the Signal, and CRT (Main Sample).

	Dependent variable: WTP ($)	
	Own strategy	Other's strategy
	(1)	(2)
$\tilde{V}$	0.534**	0.088
	(0.235)	(0.230)
DP treatment	0.176	0.085
	(0.127)	(0.134)
CRT better than partner	0.100	0.000
	(0.111)	(0.101)
CRT score	-0.009	-0.084**
	(0.051)	(0.039)
Constant	0.443***	0.670***
	(0.164)	(0.154)
Observations	67	67
R^2	0.123	0.071

Note: *p<0.1; **p<0.05; ***p<0.01

Heteroskedasticity-robust (HC1) standard errors. $\tilde{V}$: perceived value of the signal implied by the Grether parameters of the strategy that is implemented (β_i for own strategy, β_{-i} for the other's strategy), as in Section 5. Re-assigning treatments across the 12 sessions (randomization inference) gives $p = 0.333$ for the DP coefficient in column (1).

value of the signal depends strongly on the prior: it is highest at the 50% prior if a participant who does not buy the sample guesses the more likely box, and lowest at the 50% prior if the guess is random (Appendix Table D2). Bids did not follow either pattern. Average bids were \$0.55, \$0.61, and \$0.63 at the 20%, 50%, and 90% priors, and within-participant tests do not reject equal bids across priors ($p = 0.79$ for one's own strategy and $p = 0.36$ for the partner's). Only 9 of 52 participants bid the most at the 50% prior, and bids were unrelated to the value of the strategy that would be implemented at each prior. At the 50% prior, bids for one's own strategy and for the partner's strategy were similar (\$0.61 and \$0.60), as in the CP treatment.

8.5 CRT

The CRT is implemented to measure relative self-reported performance. It was expected to observe that considering themselves as more intelligent than others affects the WTP for information. To avoid experimenter demand effects, the WTP was conducted in a between-subjects design, and the effect of the CRT can only be measured at the group level and not the differences between the WTP_i and WTP_{-i} at the individual level.

Overall, 36% of the participants believed that they did better in the CRT than the participant they were paired with. On average, participants' actual CRT score was 1.06 points lower than the score they expected for themselves ($p < 0.001$), so participants were overconfident about their performance.

Table 8 shows that believing to have done better than the partner is not related to WTP. The CRT score is negatively related to WTP for the other participant's strategy (-0.08, $p = 0.04$) but not to WTP for one's own strategy (-0.01, $p = 0.86$).

9 Discussion

On average, people expect others to update in the same way they do. The dummy variable for participants' strategy or the expected strategy of others had no significant effect on the reported posterior belief. Also, the parameters estimated with Grether's model for first- and second-order beliefs were correlated. This relationship is stronger for β^l than for β^p, and no significant correlation is found between β^p and β^l. The only (marginally) significant differences were between β^l_i and β^l_{-i} in the DP treatment. With the present experiment settings, having different priors generates a difference only in the higher-order beliefs about the weight on the signal ($\beta^l_i > \beta^l_{-i}$). In general,

no systematic effect of higher-order beliefs was found. However, having different priors affects the weight expected that others put on the signal; participants expected others to be slightly more conservative.

Previous results showed no difference or weak difference between the higher-order beliefs with new information (Evdokimov and Garfagnini, 2022). However, being asked about someone with a different prior (DP treatment) increases the absolute difference between the posterior beliefs and an expectation of others being more conservative. People expect others to update similarly to themselves, also when the two beliefs are elicited from different people. Yet in the DP treatment they bid more to implement their own strategy than the other participant's strategy, while in the CP treatment the two bids are similar. A higher expectation of conservatism in others than in themselves is consistent with this gap, but the extent to which it explains the gap needs more attention.

Sensitivity to information predicts WTP. Participants who put more weight on the signal (β^l) bid more to implement their own strategy, and the value of the signal implied by their own updating predicts these bids, in the direction predicted by Grether's model. The weight on the prior (β^p) is also negatively related to bids, although the prior in the WTP task was uniform, so base-rate neglect does not change the value of the signal in that task. It might be the case that β^p correlates with other variables that affect WTP.

Taken together, these results show that sensitivity to information predicts WTP, but that two behavioral patterns should be considered. First, participants largely expect others to update as they do, so implementing their own strategy rather than the other participant's should carry little instrumental value. Yet participants in the DP treatment bid more for information when it would be used with their own strategy, while bids for the partner's strategy are the same in all sessions, and the premium does not change when the perceived value of the signal is controlled for. This pattern is consistent with a non-instrumental value of decision rights (Bartling et al., 2014; Owens et al., 2014) that becomes salient when participants are made aware that others may hold different beliefs. Because treatments were assigned by session and the difference between treatments is not statistically significant, this is suggestive rather than causal evidence, and a design that randomizes the treatment within sessions is needed to establish it. Second, WTP is not sensitive to the prior as theory predicts.

Optimal provision of information is relevant in a context where the knowledge

of the world is not perfect, and we can put some effort into finding information to make better decisions. This is the case in many settings since we don't know the world perfectly. Information provision is more relevant when the decisions of others affect our outcomes. If the incentives are aligned, it is in the agent's best interest to improve the decision of other participants. This experiment shows that the provision of information to others is hampered by people's biases.

References

Agranov, Marina and Polina Detkova, "Beliefs of Others: An Experiment," 2024. Working Paper.

Bartling, Björn, Ernst Fehr, and Holger Herz, "The Intrinsic Value of Decision Rights," *Econometrica*, 2014, *82*, 2005–2039.

Benjamin, Daniel J., "Errors in probabilistic reasoning and judgment biases," in B. Douglas Bernheim, Stefano DellaVigna, and David Laibson, eds., *Handbook of Behavioral Economics: Applications and Foundations 1*, Vol. 2, North-Holland, 2019, pp. 69–186.

_ , Aaron Bodoh-Creed, and Matthew Rabin, "Base-Rate Neglect: Foundations and Implications," 2019. Working Paper.

Blanco, Mariana, Dirk Engelmann, Alexander K Koch, and Hans-Theo Normann, "Belief elicitation in experiments: is there a hedging problem?," *Experimental Economics*, 2010, *13*, 412–438.

Brandts, Jordi and Gary Charness, "The strategy versus the direct-response method: a first survey of experimental comparisons," *Experimental Economics*, 2011, *14*, 375–398.

Campos-Mercade, Pol and Friederike Mengel, "Non-Bayesian Statistical Discrimination," *Management Science*, 2024, *70*, 2549–2567.

Charness, Gary, Ramón Cobo-Reyes, Natalia Jiménez, Juan A. Lacomba, and Francisco Lagos, "The Hidden Advantage of Delegation: Pareto Improvements in a Gift Exchange Game," *American Economic Review*, 2012, *102*, 2358–79.

_ , Uri Gneezy, and Vlastimil Rasocha, "Experimental methods: Eliciting beliefs," *Journal of Economic Behavior and Organization*, 2021, *189*, 234–256.

Danz, David, Lise Vesterlund, and Alistair J Wilson, "Belief Elicitation and Behavioral Incentive Compatibility," *American Economic Review*, 2022, *112*, 2851–2883.

Engelmann, Dirk and Martin Strobel, "The False Consensus Effect Disappears if Representative Information and Monetary Incentives Are Given," *Experimental Economics*, 2000, *3*, 241–260.

Epstein, Larry G., Jawwad Noor, and Alvaro Sandroni, "Non-Bayesian updating: a theoretical framework," *Theoretical Economics*, 2008, *3*, 193–229.

Erkal, Nisvan, Lata Gangadharan, and Han Koh, "Belief Elicitation with Binary Outcomes: A Comparison of Quadratic and Binarized Scoring Rules," *SSRN Working Paper*, 2020.

Evdokimov, Piotr and Umberto Garfagnini, "Higher-order learning," *Experimental Economics*, 2022, *25*, 1234–1266.

_ **and** _ , "Cognitive Ability and Perceived Disagreement in Learning," 2023. Discussion Paper, CRC TRR 190 Rationality and Competition.

Falk, Armin and Michael Kosfeld, "The Hidden Costs of Control," *American Economic Review*, 2006, *96*, 1611–1630.

Fedyk, Anastassia, "Asymmetric Naivete: Beliefs About Self-Control," *SSRN Electronic Journal*, 2021.

Grether, David M, "Bayes Rule as a Descriptive Model: The Representativeness Heuristic," *The Quarterly Journal of Economics*, 1980, *95*, 537–557.

_ , "Testing Bayes rule and the representativeness heuristic: Some experimental evidence," *Journal of Economic Behavior and Organization*, 1992, *17*, 31–57.

Hoppe, Eva I and David J Kusterer, "Behavioral biases and cognitive reflection," *Economics Letters*, 2011, *110*, 97–100.

Hossain, Tanjim and Ryo Okui, "The binarized scoring rule," *Review of Economic Studies*, 2013, *80*, 984–1001.

Jin, Ginger Zhe, Michael Luca, and Daniel Martin, "Is No News (Perceived As) Bad News? An Experimental Investigation of Information Disclosure," *American Economic Journal: Microeconomics*, 2021, *13*, 141–173.

Kirchkamp, Oliver, Joerg Oechssler, and Andis Sofianos, "The Binary Lottery Procedure does not induce risk neutrality in the Holt & Laury and Eckel & Grossman tasks," *Journal of Economic Behavior and Organization*, 2021, *185*, 348–369.

Oechssler, Jörg, Andreas Roider, and Patrick W Schmitz, "Cognitive abilities and behavioral biases," *Journal of Economic Behavior and Organization*, 2009, *72*, 147–152.

Owens, David, Zachary Grossman, and Ryan Fackler, "The Control Premium: A Preference for Payoff Autonomy," *American Economic Journal: Microeconomics*, 2014, *6*, 138–161.

Primi, Caterina, Kinga Morsanyi, Francesca Chiesi, Maria Anna Donati, and Jayne Hamilton, "The Development and Testing of a New Version of the Cognitive Reflection Test Applying Item Response Theory (IRT)," *Journal of Behavioral Decision Making*, 2016, *29*, 453–469.

Pronin, Emily, "How We See Ourselves and How We See Others," *Science*, 2008, *320*, 1177–1180.

_ , **Daniel Y. Lin, and Lee Ross**, "The Bias Blind Spot: Perceptions of Bias in Self Versus Others," *Personality and Social Psychology Bulletin*, 2002, *28*, 369–381.

Schotter, Andrew and Isabel Trevino, "Belief elicitation in the laboratory," *Annual Review of Economics*, 2014, *6*, 103–128.

Szkup, Michal and Isabel Trevino, "Sentiments, strategic uncertainty, and information structures in coordination games," *Games and Economic Behavior*, 2020, *124*, 534–553.

_ **and** _ , "Selection through information acquisition in coordination games," *Economic Theory*, 2026, *81* (1), 113–148.

A Instructions: Payoff Calculations for Section 1

In every task, the computer will select a Box, generate a sample of 4 marbles, and implement your guess about the probability of the Yellow Box selected (p) to calculate your payoffs. An independent Box and sample will be selected for Participant 2, who chooses the strategy q. Remember you also have to report your guess about q, and let's call it r. These two guesses you make (p and r) determine the probability of making up to $3 in each task:

1. Guess on the box selected by the computer: p (your report divided by 100)

   For every possible sample, indicate a probability p between 0 and 1 that you believe the Yellow box was selected. If the computer selects the Yellow Box, you will make $1.50 with probability $1-(1-p)^2$. If the computer selects the Green Box, you will make $1.50 with probability $1-p^2$. This is also shown below:

Selected Box	**Probability of winning**
Yellow Box	1 - $(1-p)^2$
Green Box	1- p^2

   Here are some examples of how the payment works:

   - If you think that the Yellow box was selected for sure, choosing p=1 gives your best chance of making $1.50. This happens because the probability of winning would be $1-(1-1)^2 = 1$, or 100%. Choosing any other number will decrease your chances of winning. For example, if you chose p = 0.5, the probability of winning would be $1-(1-0.5)^2 = 1-(0.5)^2 = 1-0.25 = 0.75$, or 75%; if you chose p = 0, the probability of winning would be $1-(1-0)^2 = 1-1 = 0$, or 0%.
   - Similarly, if you think that the Green box was selected for sure, choosing p=0 gives your best chance of winning; the probability of winning would be $1-(0)^2 = 1$, or 100%. If you chose p = 0.5, the probability of winning would be $1-(0.5)^2 = 1-0.25 = 0.75$, or 75%; if you chose p = 1, the probability of winning would be $1-(1)^2 = 1-1 = 0$, or 0%.
   - If you believe that the Yellow box was selected with a 50% chance, choosing p=0.5 gives your best chance of making $1.50; the chance of winning would be the average between the probabilities of winning given each box: $0.50*$

$(1-(1-0.5)^2)+0.50*(1-(0.5)^2) = 0.50*(1-0.5^2)+0.50*(1-0.5^2) = 0.50*0.75+0.50*0.75 = 0.75$, or 75%. Choosing any other number will decrease your chances. For example, if you chose p = 1, the probability of winning would be: $0.50*(1-(1-1)^2)+0.50*(1-(1)^2) = 0.50*(1)+0.50*(0) = 0.50$, or 50%. Similarly, if you chose p = 0, the probability of winning would be $0.50*(0)+0.50*(1) = 0.5$, or 50%.

Any number different from your belief that Yellow Box was selected will decrease your chances of winning.

2. Guess on Participant 2's guess: r (your report divided by 100)

   Participant 2 will also bet on the probability that a Yellow box was selected for them by reporting the number q. You will indicate a probability r between 0 and 1, which is your guess for q. You will make $1.5 with probability:

$$(1-(q-r)^2)$$

   You would wish to choose r to be as close as possible to q. Participant 2 will have the same payment scheme as you do in your first guess above. Here are some examples of how the payment works:

   - If you think that Participant 2 guessed the Yellow box was selected for sure (q=1), choosing r=1 maximizes your chances of making $1.50. The probability of winning would be $1-(1-1)^2 = 1$, or 100%. Choosing any other number will decrease your chances of winning. For example, if you chose r = 0.5, the probability of winning would be $1-(1-0.5)^2 = 1-0.5^2 = 0.75$, or 75%; if you chose r = 0, the probability of winning would be $1-(1-0)^2 = 0$, or 0%.
   - If you think that Participant 2 guessed the Yellow box was selected with a 50% probability (q=0.5), choosing r=0.5 maximizes your chances of making $1.50. The probability of winning would be $1-(0.5-0.5)^2 = 1-0 = 1$, or 100%. Choosing any other number will decrease your chances. For example, if you chose r = 1, the probability of winning would be $(1-(0.5-1)^2) = 1-0.25 = 0.75$ or 75%; if you chose r = 0, the probability of winning would be $(1-(0.5-0)^2) = 1-0.25 = 0.75$, or 75%.

Any number different from your belief about Participant 2's strategy will decrease your chances of winning.

These bets were designed such that truthfully reporting your beliefs maximizes your chances of getting up to $9 in this section. Be as accurate as possible, and carefully consider the probability of the Yellow Box selected given each sample and what Participant 2 would choose.

B Instructions: Payoff Calculations for Section 2

The computer will select the Yellow or Green Box with the same probability: 50%. That means that a random Box will be selected from Cabinet 2.

In this section, you and Participant 2 will be paid $3 if the decision about the box selected by the computer is correct. The strategies in Section 1 will determine the decision to make considering a sample generated from the selected box:

- If the probability reported was more than 50% for the Yellow Box, the decision will be to choose the Yellow Box.
- If the probability reported was less than 50% for the Yellow Box, the decision will be to choose the Green Box.
- If the probability reported was equal to 50% for the Yellow Box, the computer will randomly choose one box for you.

Willingness to Pay to Apply Participant 2's Strategy The computer will select the Yellow Box with a 50% probability and generate a sample of 4 marbles. However, in this section, you will not observe the sample for free. You have to decide how much you are willing to pay so that Participant 2's strategy reported in Section 1 can be applied to the sample generated from the Selected Box. Notice that a sample can be informative of the actual box selected.

How is your purchase determined? You will participate in an auction against the computer to purchase the right to observe the sample generated. You have to make a bet that represents your willingness to pay.

An amount of $X between $0.00 and $1.50 will be randomly chosen by the computer. $X can take any value between $0.00 and $1.50 with the same probability. Then two cases can occur concerning your bet:

If your willingness to pay is larger or equal to $X, you will pay $X to observe the sample. Notice that $X cannot be larger than your bet in this case. If your willingness to pay is smaller than $X, you pay nothing, and no sample is shown. Notice that if the sample does not increase your chances of selecting the correct box and making $3, you should bet $0.00. And if the sample increases your chances, you should increase your bet. You will increase your expected payoffs if you report how much extra money

you expect to make if Participant 2's strategy in Section 1 is applied to the sample observed.

Remember that you and Participant 2 will be paid \$3 if Participant 2's decision about the box selected by the computer is correct.

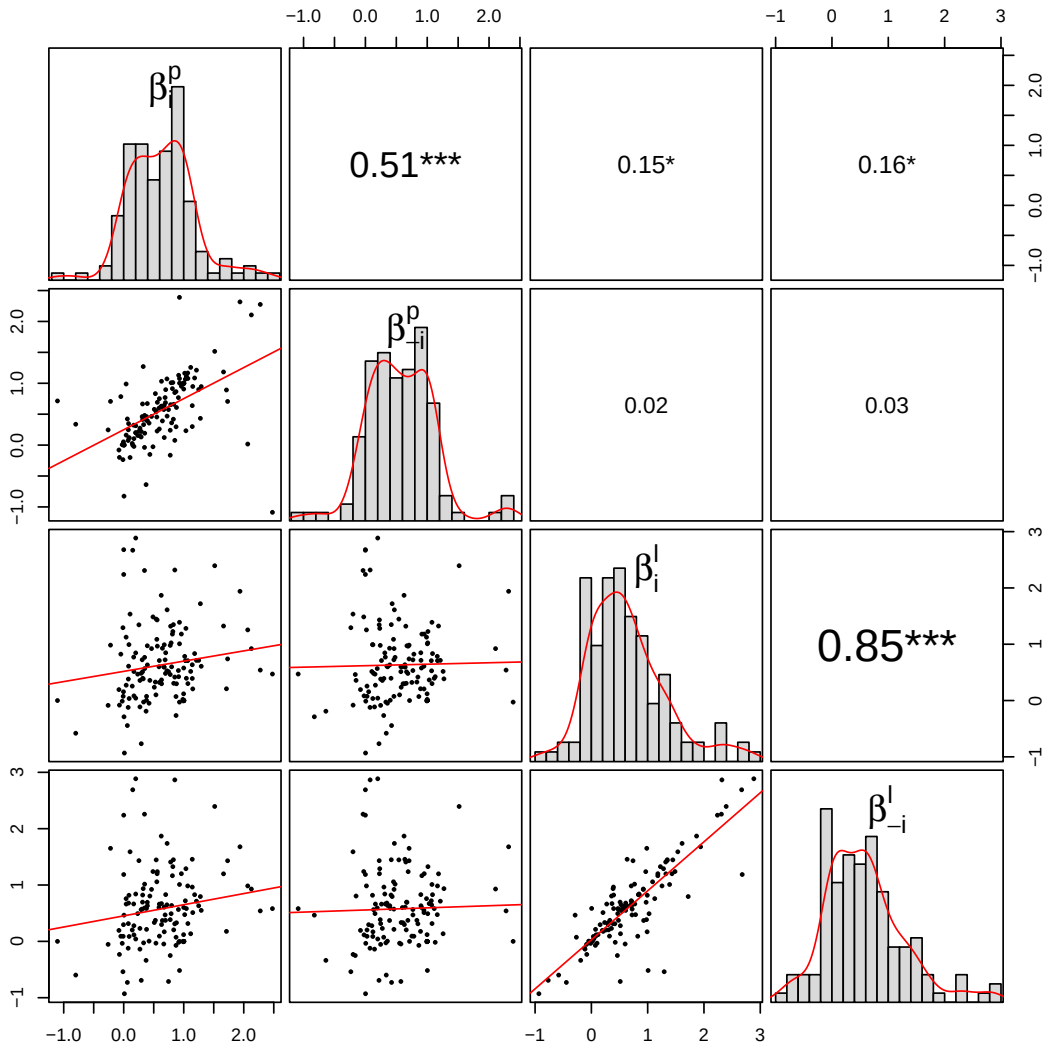
C Individual Grether Coefficients

Appendix Figures C1 and C2 show the correlations of the individually estimated Grether coefficients for all participants and for the DP treatment only. Appendix Table C1 reports additional regressions of WTP on the individual coefficients.

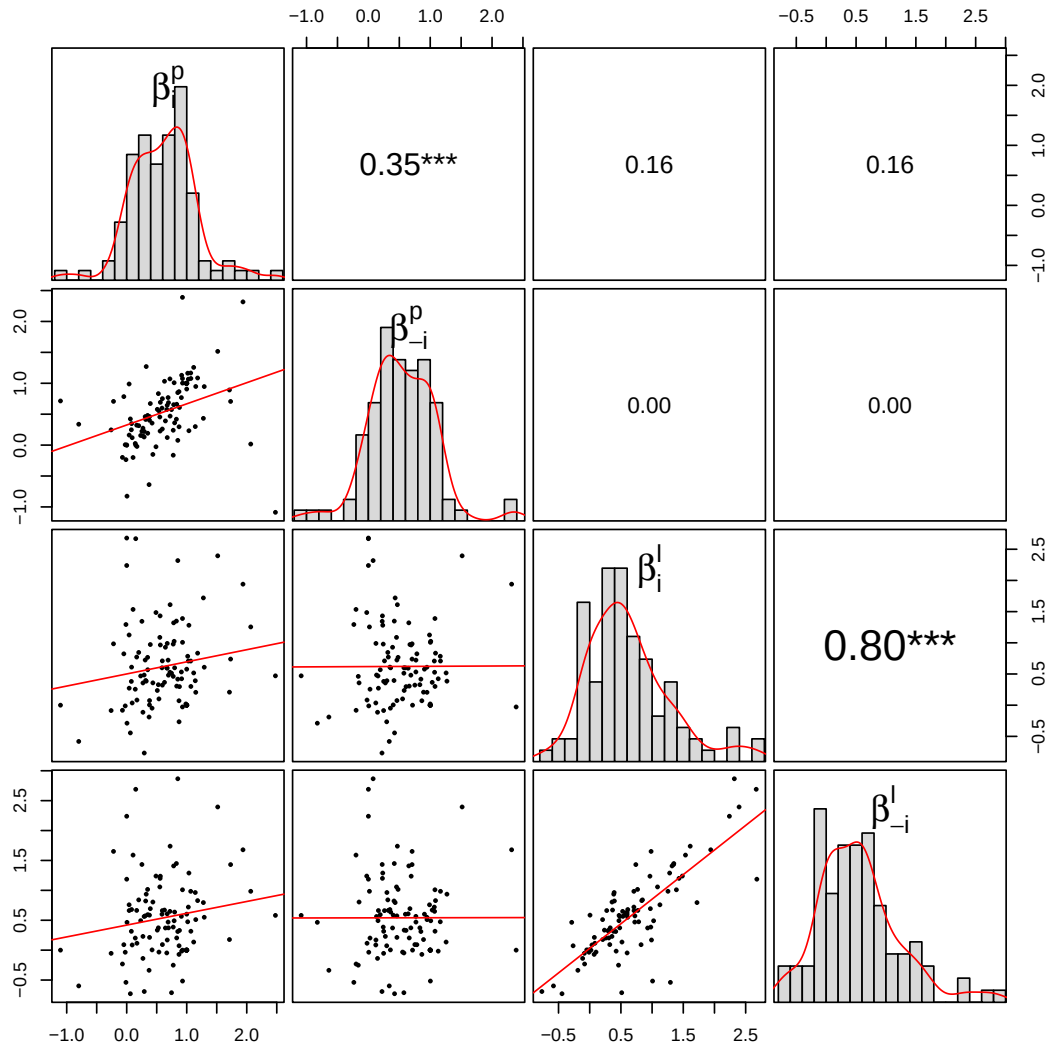
Appendix Table C1. Regression of the WTP on First- and Second-Order Beliefs about β^l.

	Dependent variable: WTP ($)		
	(1)	(2)	(3)
Own strategy (integrated)	0.117	0.107	0.130*
	(0.073)	(0.072)	(0.071)
CRT better than partner	0.039		
	(0.076)		
CRT score	-0.045		
	(0.031)		
β^p (pooled)	-0.195***	-0.220***	-0.202***
	(0.070)	(0.070)	(0.072)
β^l_{-i}	-0.017	-0.028	
	(0.094)	(0.096)	
β^l_i	0.131	0.140	
	(0.096)	(0.097)	
CP treatment	-0.093		
	(0.088)		
$\beta^l_i - \beta^l_{-i}$			0.072
			(0.098)
Constant	0.772***	0.712***	0.758***
	(0.082)	(0.063)	(0.060)
Observations	134	134	134
R^2	0.139	0.115	0.086

Note: *p<0.1; **p<0.05; ***p<0.01
Heteroskedasticity-robust (HC1) standard errors. β^p (pooled) is the prior coefficient estimated from all of a participant's reports (own and predicted-other).



Appendix Figure C1. Correlation of the coefficients from Grether's model estimated individually. This includes the DP and CP treatments.



Appendix Figure C2. Correlation of the coefficients from Grether's model estimated individually. This includes only the DP treatment.

D Additional Sessions

This appendix reports the two sets of additional sessions described in Sections 7.3 and 8. Neither is part of the main sample. Appendix Table D1 compares the between-subject sessions with the two treatments of the main sample, and Appendix Table D2 reports their bids at each prior. Appendix Table D3 compares the Common Priors sessions with the monotonicity check with the CP treatment of the main sample.

Appendix Table D1. Between-Subject Elicitation Compared with the Main Sample.

Measure	BS	CP (main)	DP (main)	p: BS vs CP	p: BS vs DP
WTP at the 50% prior, own strategy	0.61	0.64	0.82	0.821	0.061
WTP at the 50% prior, other's strategy	0.60	0.58	0.64	0.902	0.658
Mean |own posterior − Bayes| (p.p.)	19.15	20.01	19.13	0.729	0.994
Mean |predicted-other posterior − Bayes| (p.p.)	19.64	20.22	20.39	0.842	0.758
Median β_i^l (own)	0.48	0.59	0.51	0.600	0.431
Median β_i^p (own)	0.51	0.63	0.64	0.260	0.281
Median β_{-i}^l (predicted other)	0.71	0.63	0.49	0.681	0.098
Median β_{-i}^p (predicted other)	0.50	0.63	0.47	0.810	0.672

BS: between-subject sessions (52 participants; common priors), in which each participant reported either their own posteriors or their prediction of the partner's posteriors. CP and DP: treatments of the main sample. WTP p-values: Welch t-tests. Posterior p-values: regressions with standard errors clustered by participant. Medians: Mann-Whitney tests.

Appendix Table D2. Willingness to Pay at Each Prior in the Between-Subject Sessions.

Prior	Own strategy	Other's strategy	All	Bayesian WTP (no purchase: likelier box)	Bayesian WTP (no purchase: random guess)
20%	0.615	0.479	0.547	0.016	0.916
50%	0.612	0.602	0.607	0.444	0.444
90%	0.612	0.656	0.634	0.000	1.200

Mean bids ($); 26 participants per strategy, each bidding once at each prior. Bayesian benchmarks are shown for the two possible rules when no sample is purchased. Friedman tests of equal bids across priors: own strategy p = 0.786, other's strategy p = 0.364, all p = 0.649.

Appendix Table D3. Common Priors Sessions with and without the Monotonicity Check.

Measure	CP 2023	CP 2024 (monotonicity check)	p-value
Mean WTP, own strategy (WTP_i)	0.64	0.67	0.840
Mean WTP, other's strategy (WTP_{-i})	0.58	0.64	0.719
Own − other's WTP gap	0.06	0.03	0.898
Share of strategy tables monotone in the signal	0.72	1.00	<0.001
Mean |posterior − Bayes| (p.p.)	20.12	17.04	0.127
Median β_i^l (own)	0.59	0.95	0.004
Median β_{-i}^l (predicted other)	0.63	1.09	0.001
Median β_i^p (own)	0.63	0.52	0.488
Median β_{-i}^p (predicted other)	0.63	0.44	0.288
Share of cells where own = predicted-other posterior	0.74	0.58	<0.001
Mean |own − predicted-other posterior| (p.p.)	2.67	4.62	0.051

CP 2023: Common Priors sessions of the main sample (32 participants). CP 2024: Common Priors sessions run in May 2024 (56 participants) with a check that flagged strategy tables that were not monotone in the number of yellow marbles. WTP p-values: Welch t-tests (gap: robust test of the year × own-strategy interaction). Posterior p-values: regressions with standard errors clustered by participant; proportions: χ^2 tests; medians: Mann-Whitney tests.